

# Progressive Neural Networks

Trofimova Yuliya

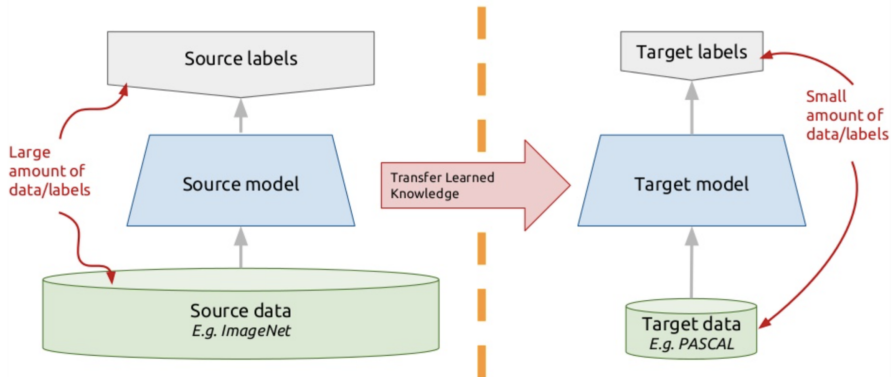
Higher School of Economics

19 декабря 2019 г.


# Overview

- 1 Transfer learning
- 2 Progressive Network
- 3 Adapter
- 4 Transfer Analysis
- 5 Experiments
- 6 Experiment1: Pong Soup
- 7 Experiment2: Atari
- 8 Experiment3: Labyrinth

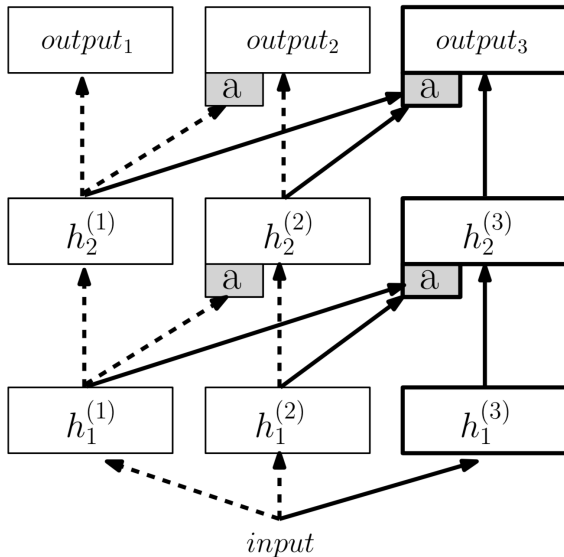
# Transfer learning



# Finetuning



# Progressive Network



$$h_i^{(k)} = f \left( W_i^{(k)} h_{i-1}^{(k)} + \sum_{j < k} U_i^{(k:j)} h_{i-1}^{(j)} \right)$$

- $L$  layers
- $h_i^{(k)} \in \mathbb{R}^{n_i}$ , with  $n_i$  - number of units at layer  $i \leq L$
- $\theta^{(k-1)}$  are "frozen"
- $W_i^{(k)} \in \mathbb{R}^{n_i \times n_{i-1}}$  - weight matrix of layer  $i$  of column  $k$
- $U_i^{k:j} \in \mathbb{R}^{n_i \times n_j}$  - lateral connections from layer  $i-1$  of column  $j$ , to layer  $i$  of column  $k$  and  $h_0$  is the network input
- $f(x) = \max(0, x)$  - element-wise non-linearity

Define vector of anterior features of dimensionality  $n_{i-1}^{(<k)}$

$$h_{i-1}^{(<k)} = [h_{i-1}^{(1)} \cdots h_{i-1}^{(j)} \cdots h_{i-1}^{(k-1)}]$$

$$h_i^{(k)} = \sigma \left( W_i^{(k)} h_{i-1}^{(k)} + U_i^{(k:j)} \sigma(V_i^{(k:j)} \alpha_{i-1}^{(<k)} h_{i-1}^{(<k)}) \right)$$

- $\alpha_{i-1}^{(<k)}$  - random small value
- $V_i^{(k:j)}$  - projection onto an  $n_i$  dimensional subspace

When  $k$  grows, number of parameters in lateral connections is in the same as  $\Theta^{(1)}$

# Transfer Analysis

Unlike finetuning, progressive nets do not destroy the features learned on prior tasks.

- **Average Perturbation Sensitivity:** inject Gaussian noise at isolated points in the architecture (e.g. a given layer of a single column)
- **Average Fisher Sensitivity**

$$\hat{F}_i^{(k)} = \mathbb{E}_{\rho(s,a)} \left[ \frac{\partial \log \pi}{\partial \hat{h}_i^{(k)}} \frac{\partial \log \pi^T}{\partial \hat{h}_i^{(k)}} \right] \quad \text{AFS}(i, k, m) = \frac{\hat{F}_i^{(k)}(m, m)}{\sum_k \hat{F}_i^{(k)}(m, m)}$$

- $\hat{F}$  - diagonal Fisher matrix of policy  $\pi$  with respect to the normalized activations at each layer  $\hat{h}_i^{(k)}$
- $\rho(s, a)$  - state-action distribution from target task




# Reinforcement learning

The  $k$ -th column defines a policy  $\pi^{(k)}(a|s)$  taking as input a state  $s$  given by the environment, and generating probabilities over actions

$$\pi^{(k)}(a|s) = h_L^{(k)}(s)$$



(a) Pong variants




(b) Labyrinth games






(c) Atari games

# Experiments

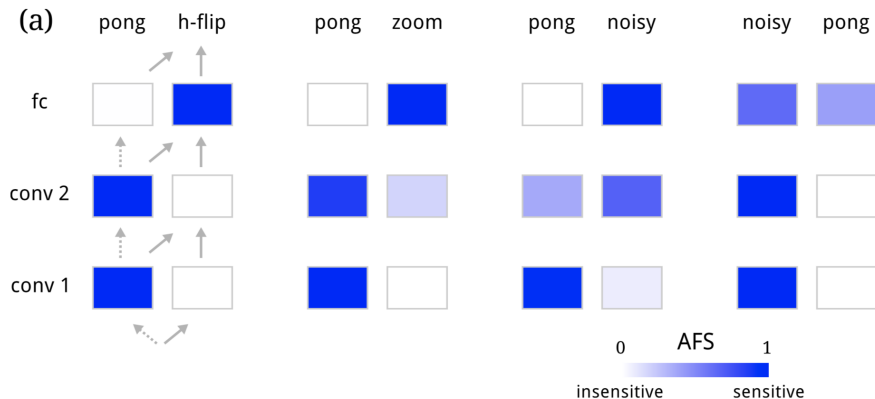


# Pong Soup

**Noisy** (frozen Gaussian noise), **Black** (black background), **White** (white background), **Zoom** (input is scaled by 75% and translated), **V-flip**, **H-flip**, and **VH-flip** (input is horizontally and/or vertically flipped)


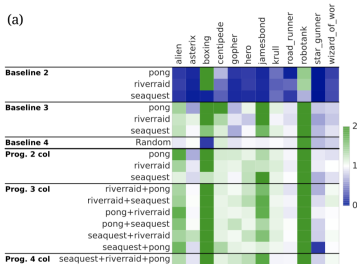


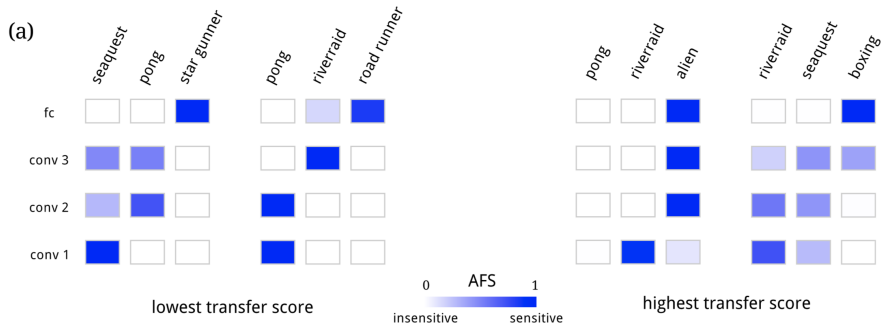
# Pong Soup AFS



Conv1 - mid-vision, Conv2 - low-vision, fc - policy

**Progressive nets** result in positive transfer in 8 out of 12 target tasks and negative transfer is 2. **Baseline 3** in positive transfer with only 5 of 12.








- The most negative transfer coincides with complete dependence on the convolutional layers of the previous columns, and no learning of new visual features in the new column.
- The most positive transfer occurs when the features of the first two columns are augmented by new features

# Labyrinth

(a)



- Seek Track 1: simple corridor with many apples
- Seek Track 2: U-shaped corridor with many strawberries
- Seek Track 3:  $\Omega$ -shaped, with 90 turns, with few apples
- Seek Track 4:  $\Omega$ -shaped, with 45 turns, with few apples
- Seek Avoid 1: large square room with apples and lemons
- Seek Avoid 2: large square room with apples and mushrooms
- Seek Maze M: M-shaped maze, with apples at dead-ends
- Seek Maze Y: Y-shaped maze, with apples at dead-ends

- Progressive neural networks are a stepping stone towards continual learning
- Progressive approach is able to effectively exploit transfer for compatible source and task domains



- Напишите формулу скрытого слоя для progressive network.
- Почему Progressive network лучше, чем finetuning всей модели, справляются с трансфером из ортогональной задачи?
- Какие зависимости наблюдаются между source и target task в случаях наибольшего негативного и позитивного переноса?