

Deep Image Prior

Кнунянц Иван БПМИ171

Задачи image restoration

Denoising



Corrupted



Deep image prior

Задачи image restoration

Inpainting



Corrupted



Deep image prior

Задачи image restoration

Super-resolution

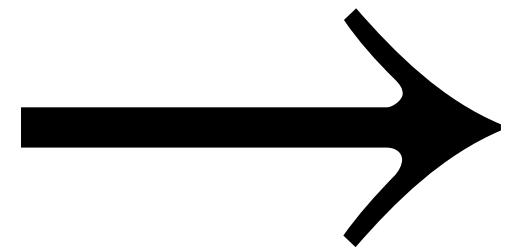


Corrupted



Deep image prior

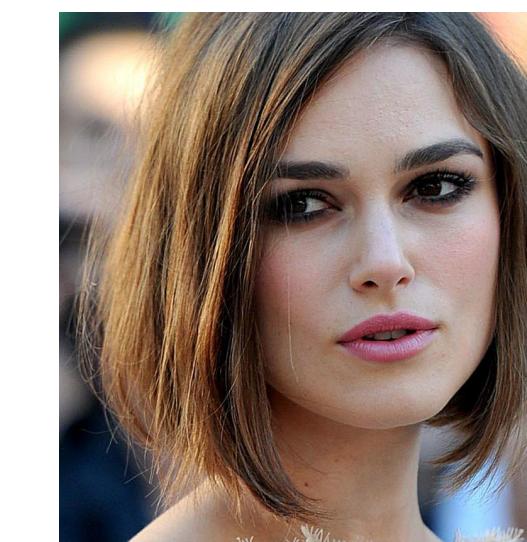
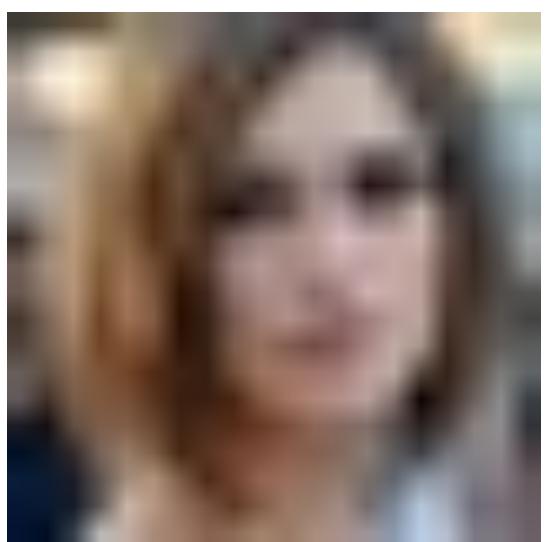
Что такое Prior?



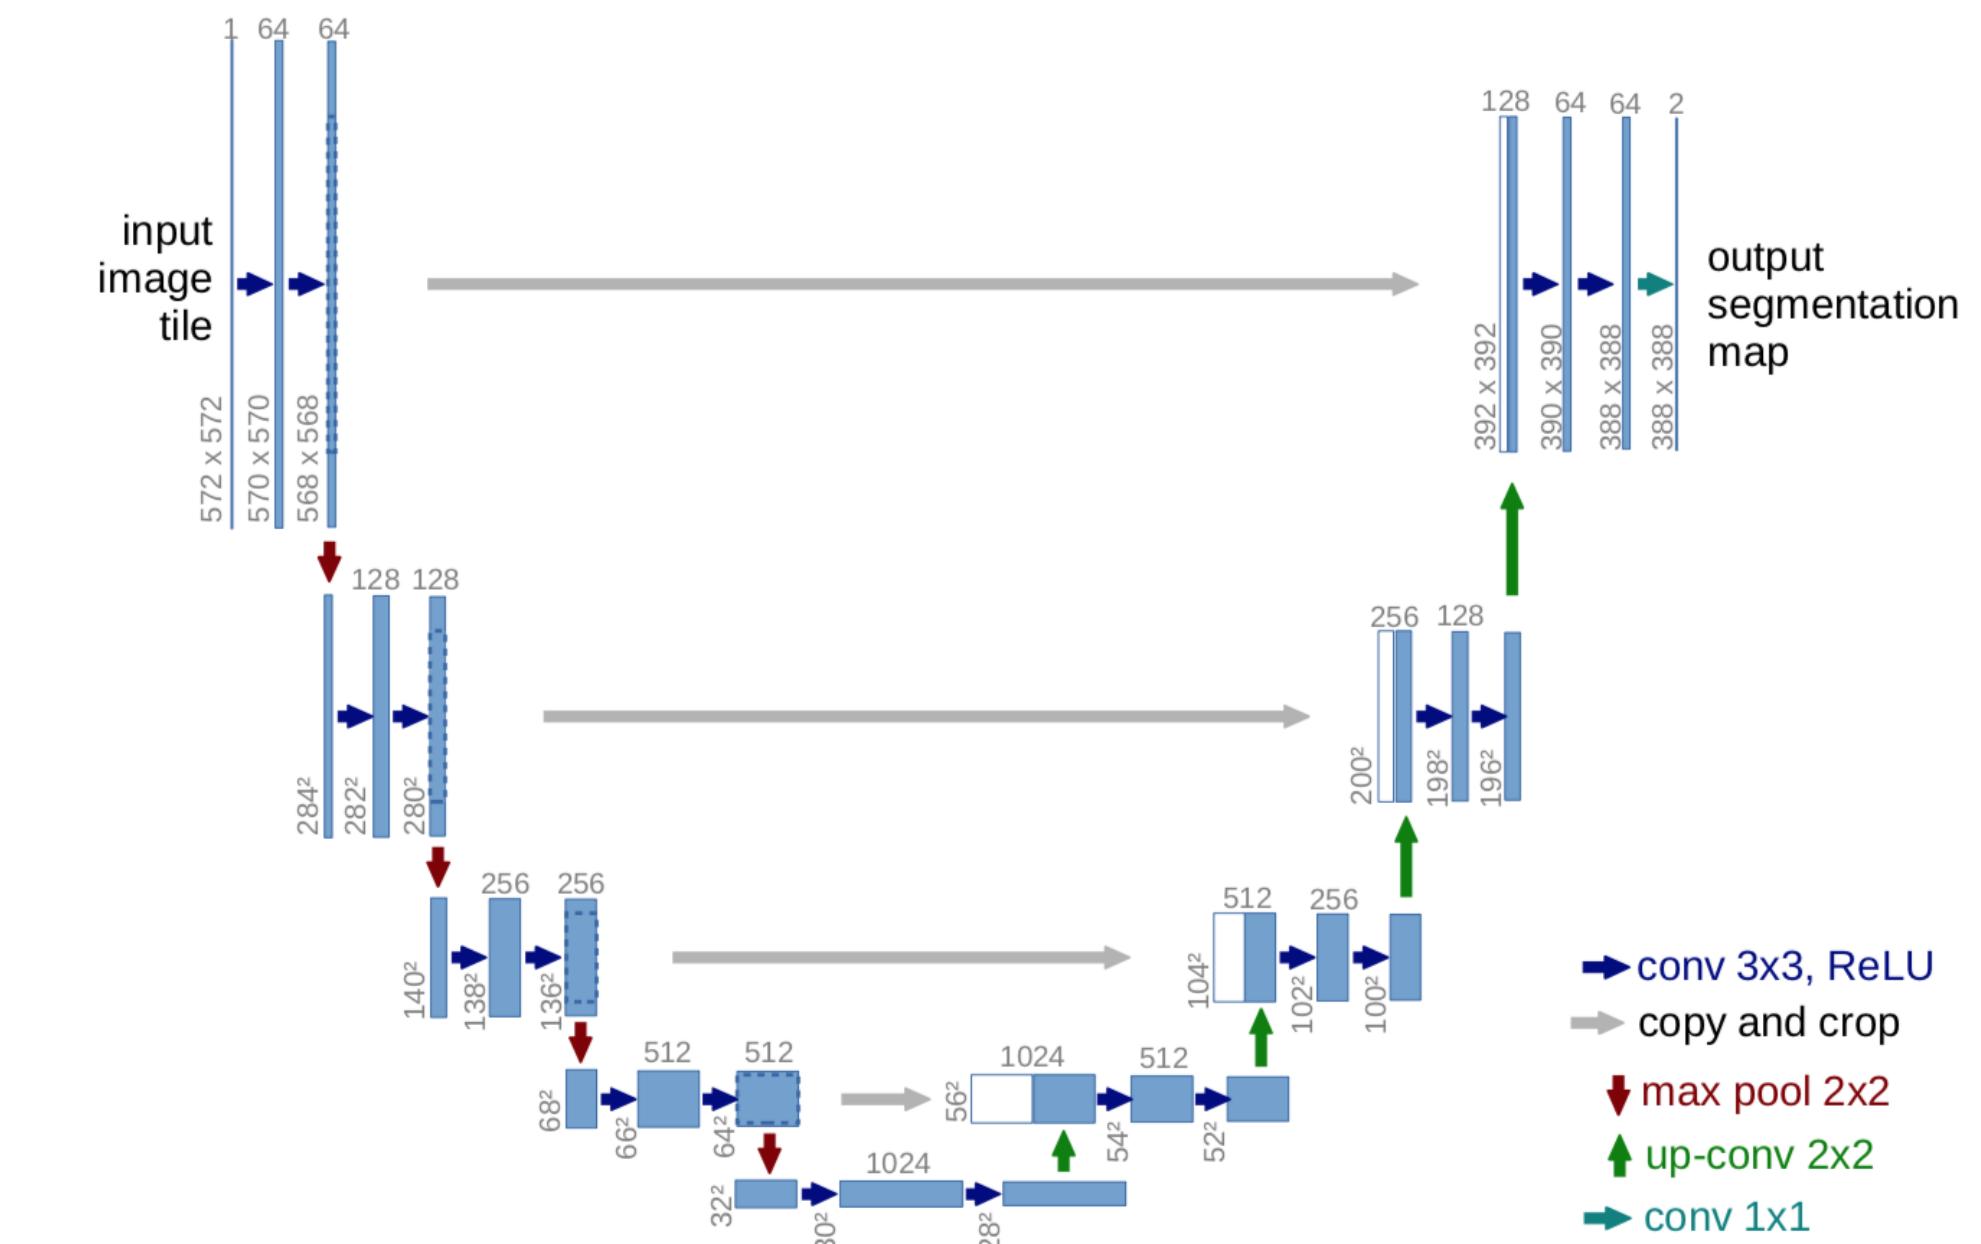
Prior - это наше “знание о мире”, которое выучивает нейросеть.

Составляющие prior'a

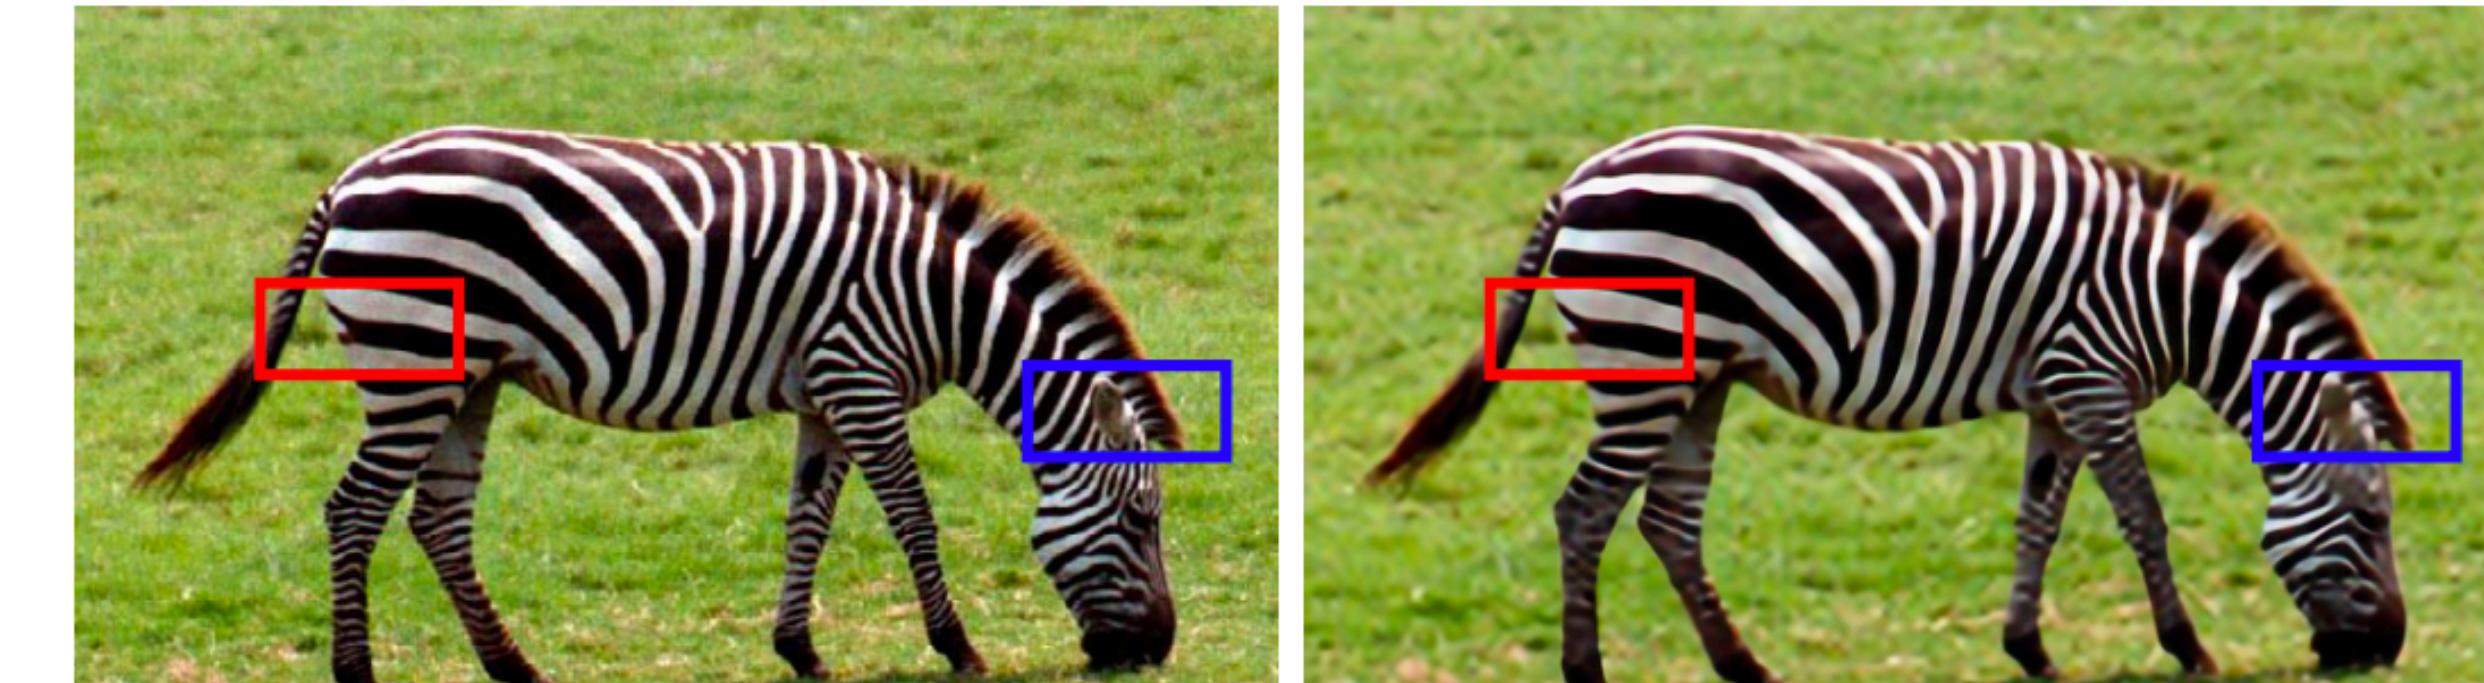
Выученный prior



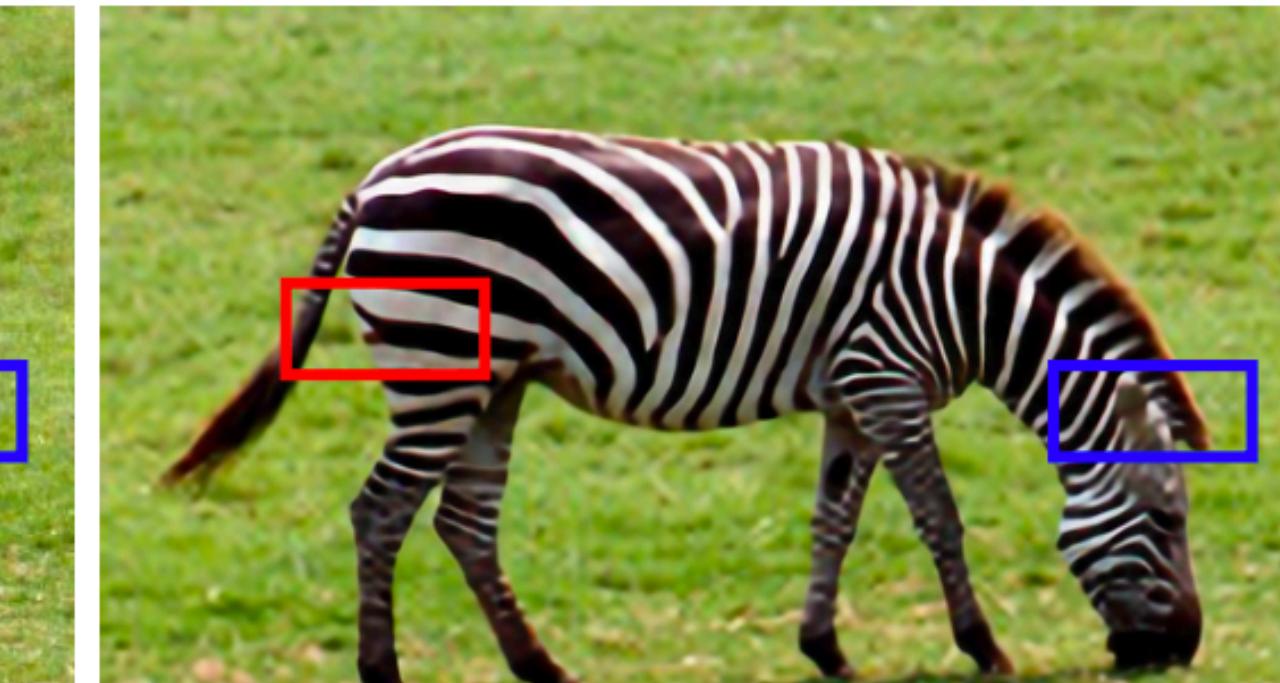
Структура нейросети



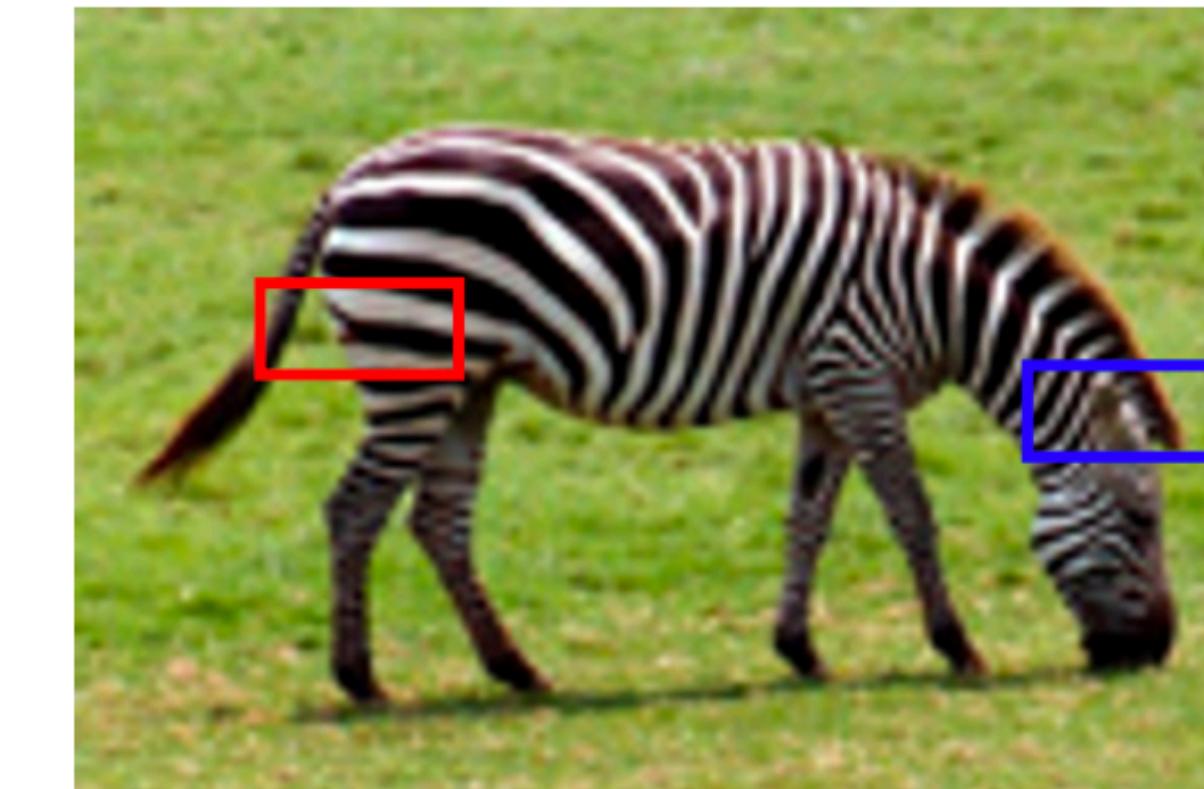
Результаты



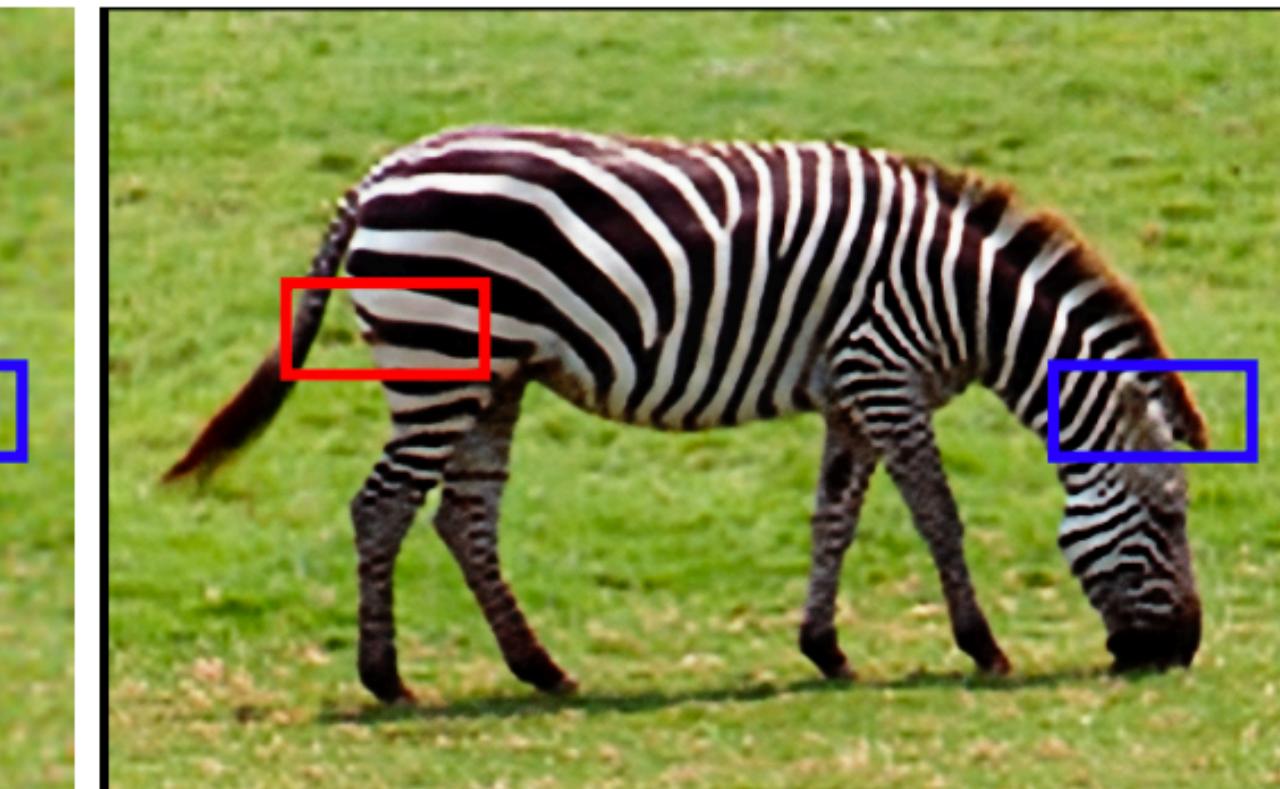
(a) Ground truth



(b) SRResNet [19], Trained



(c) Bicubic, Not trained



(d) Deep prior, Not trained

Идея image restoration

x - искомое изображение

\hat{x} - испорченное изображение

Полученное изображение: $x^* = \arg \max_x p(x | \hat{x})$

$$p(x | \hat{x}) = \frac{p(\hat{x} | x)p(x)}{p(\hat{x})} \Rightarrow p(\hat{x} | x)p(x) \rightarrow \text{maximize}$$

$p(\hat{x} | x)$ - правдоподобие, $p(x)$ - prior

Идея image restoration

Запишем выражение по-другому:

$$\begin{aligned}x^* &= \arg \max_x p(x | \hat{x}) = \arg \max_x p(\hat{x} | x)p(x) = \\&= \arg \min_x -\log(p(\hat{x} | x)) - \log(p(x)) = \arg \min_x E(x; \hat{x}) + R(x)\end{aligned}$$

$E(x; \hat{x})$ - data term, $R(x)$ - image prior

Параметризация

Пусть $g : \theta \mapsto x$ сюръективна. Тогда задача эквивалентна:

$$\arg \min_{\theta} E(g(\theta); \hat{x}) + R(g(\theta))$$

Теперь будем настраивать g , который даёт лучший результат.

Будем использовать g в качестве prior. Тогда будем оптимизировать только $\arg \min_{\theta} E(g(\theta); \hat{x})$

Зафиксируем $g \equiv f_{\theta}(z)$, где z - фиксированная величина (картина), f_{θ} - свёрточная сеть с параметрами θ

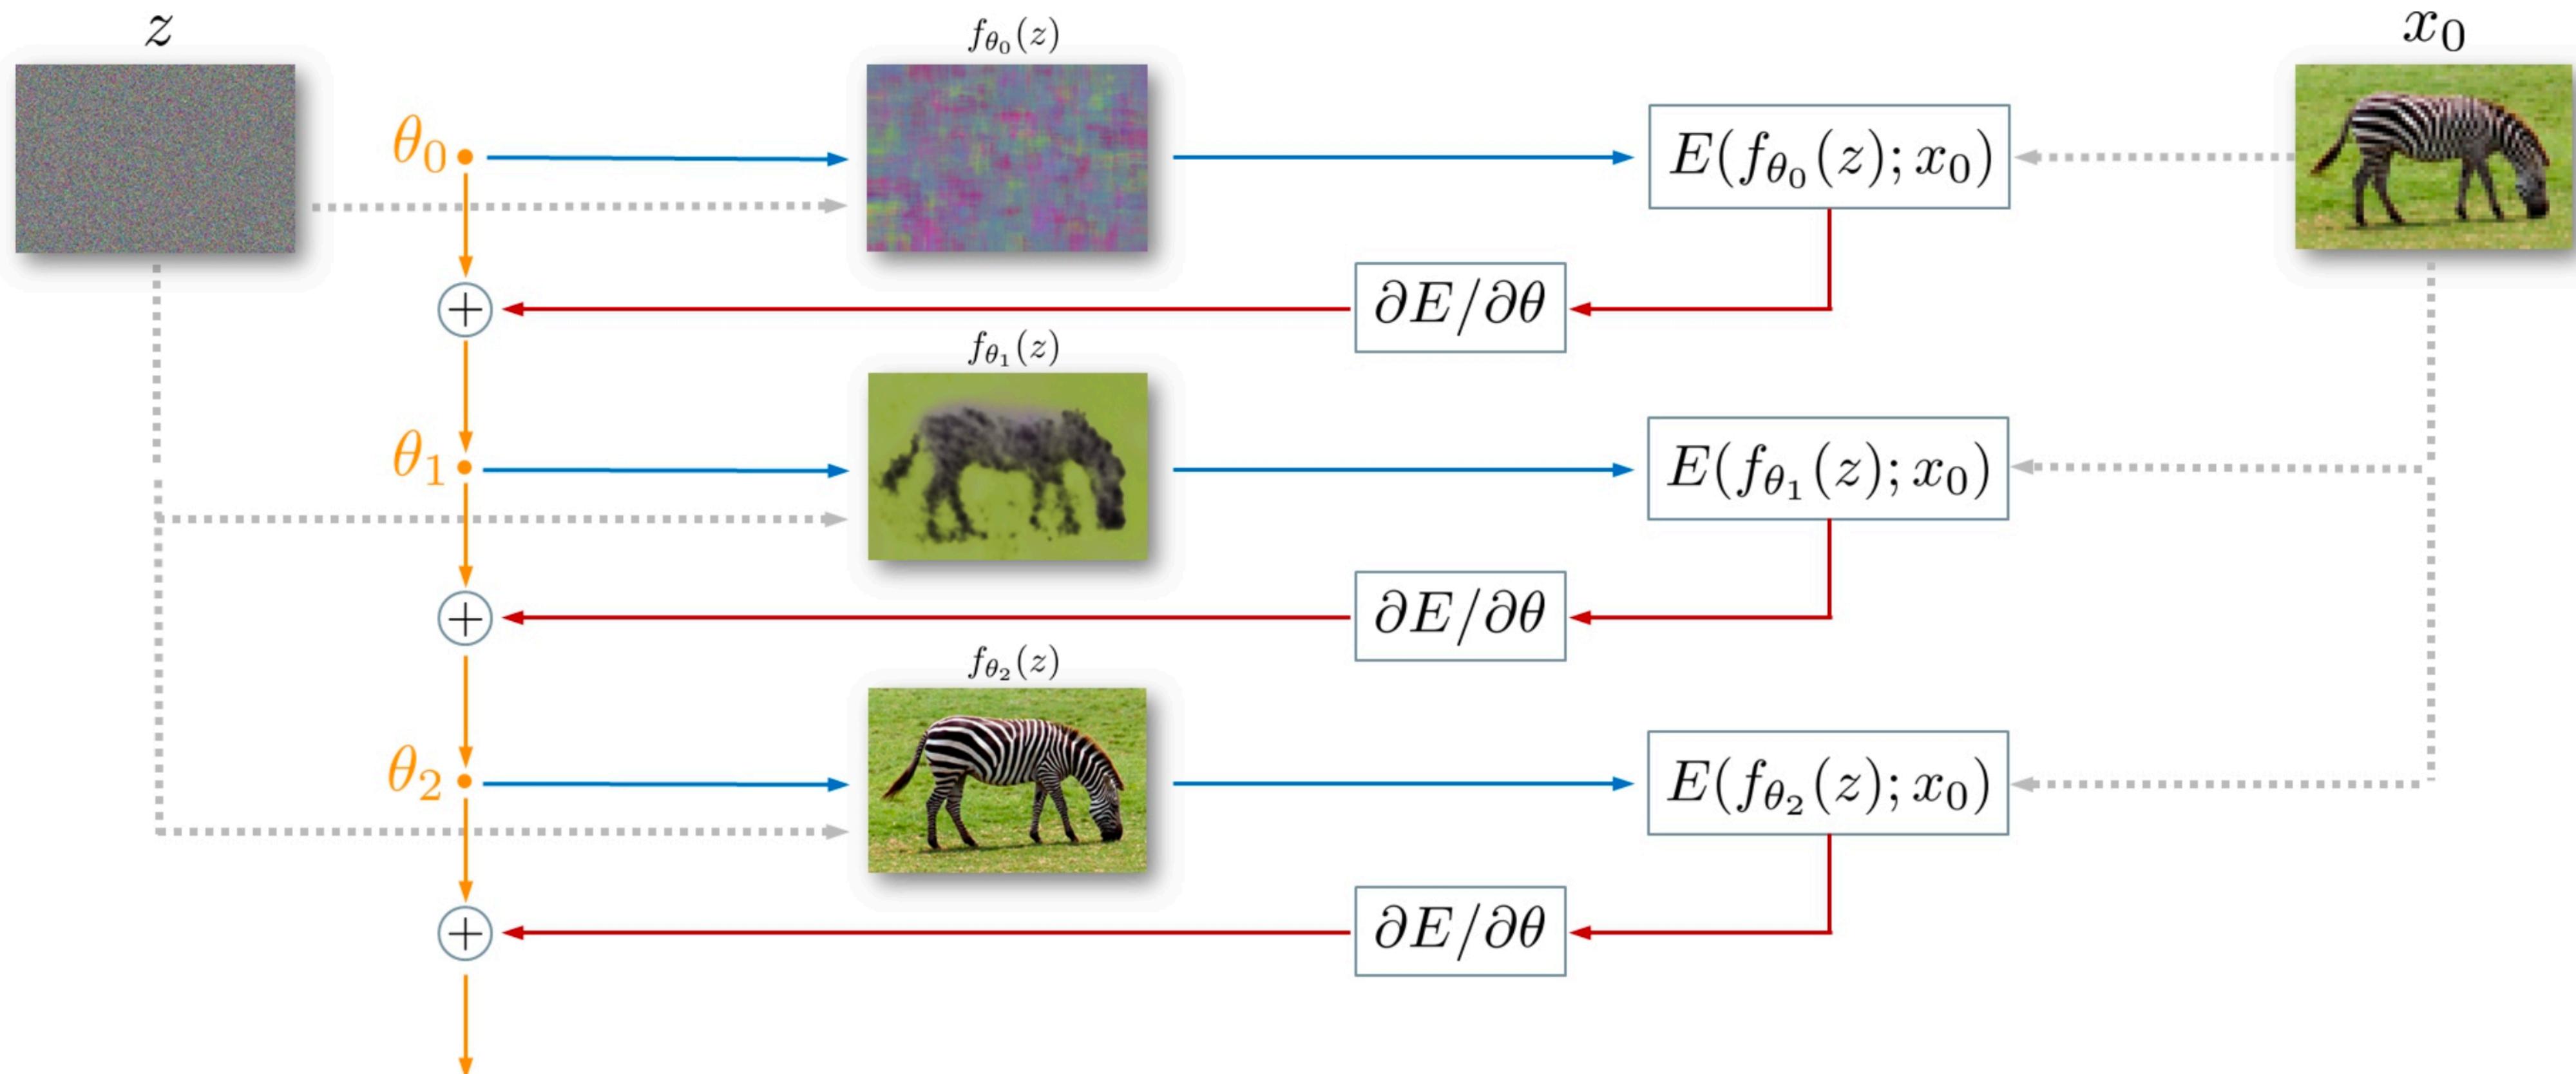
Алгоритм

1. Инициализируем z . Например, равномерным шумом $U(-1,1)$.
2. Решаем $\arg \min_{\theta} E(f_{\theta}(z); \hat{x})$

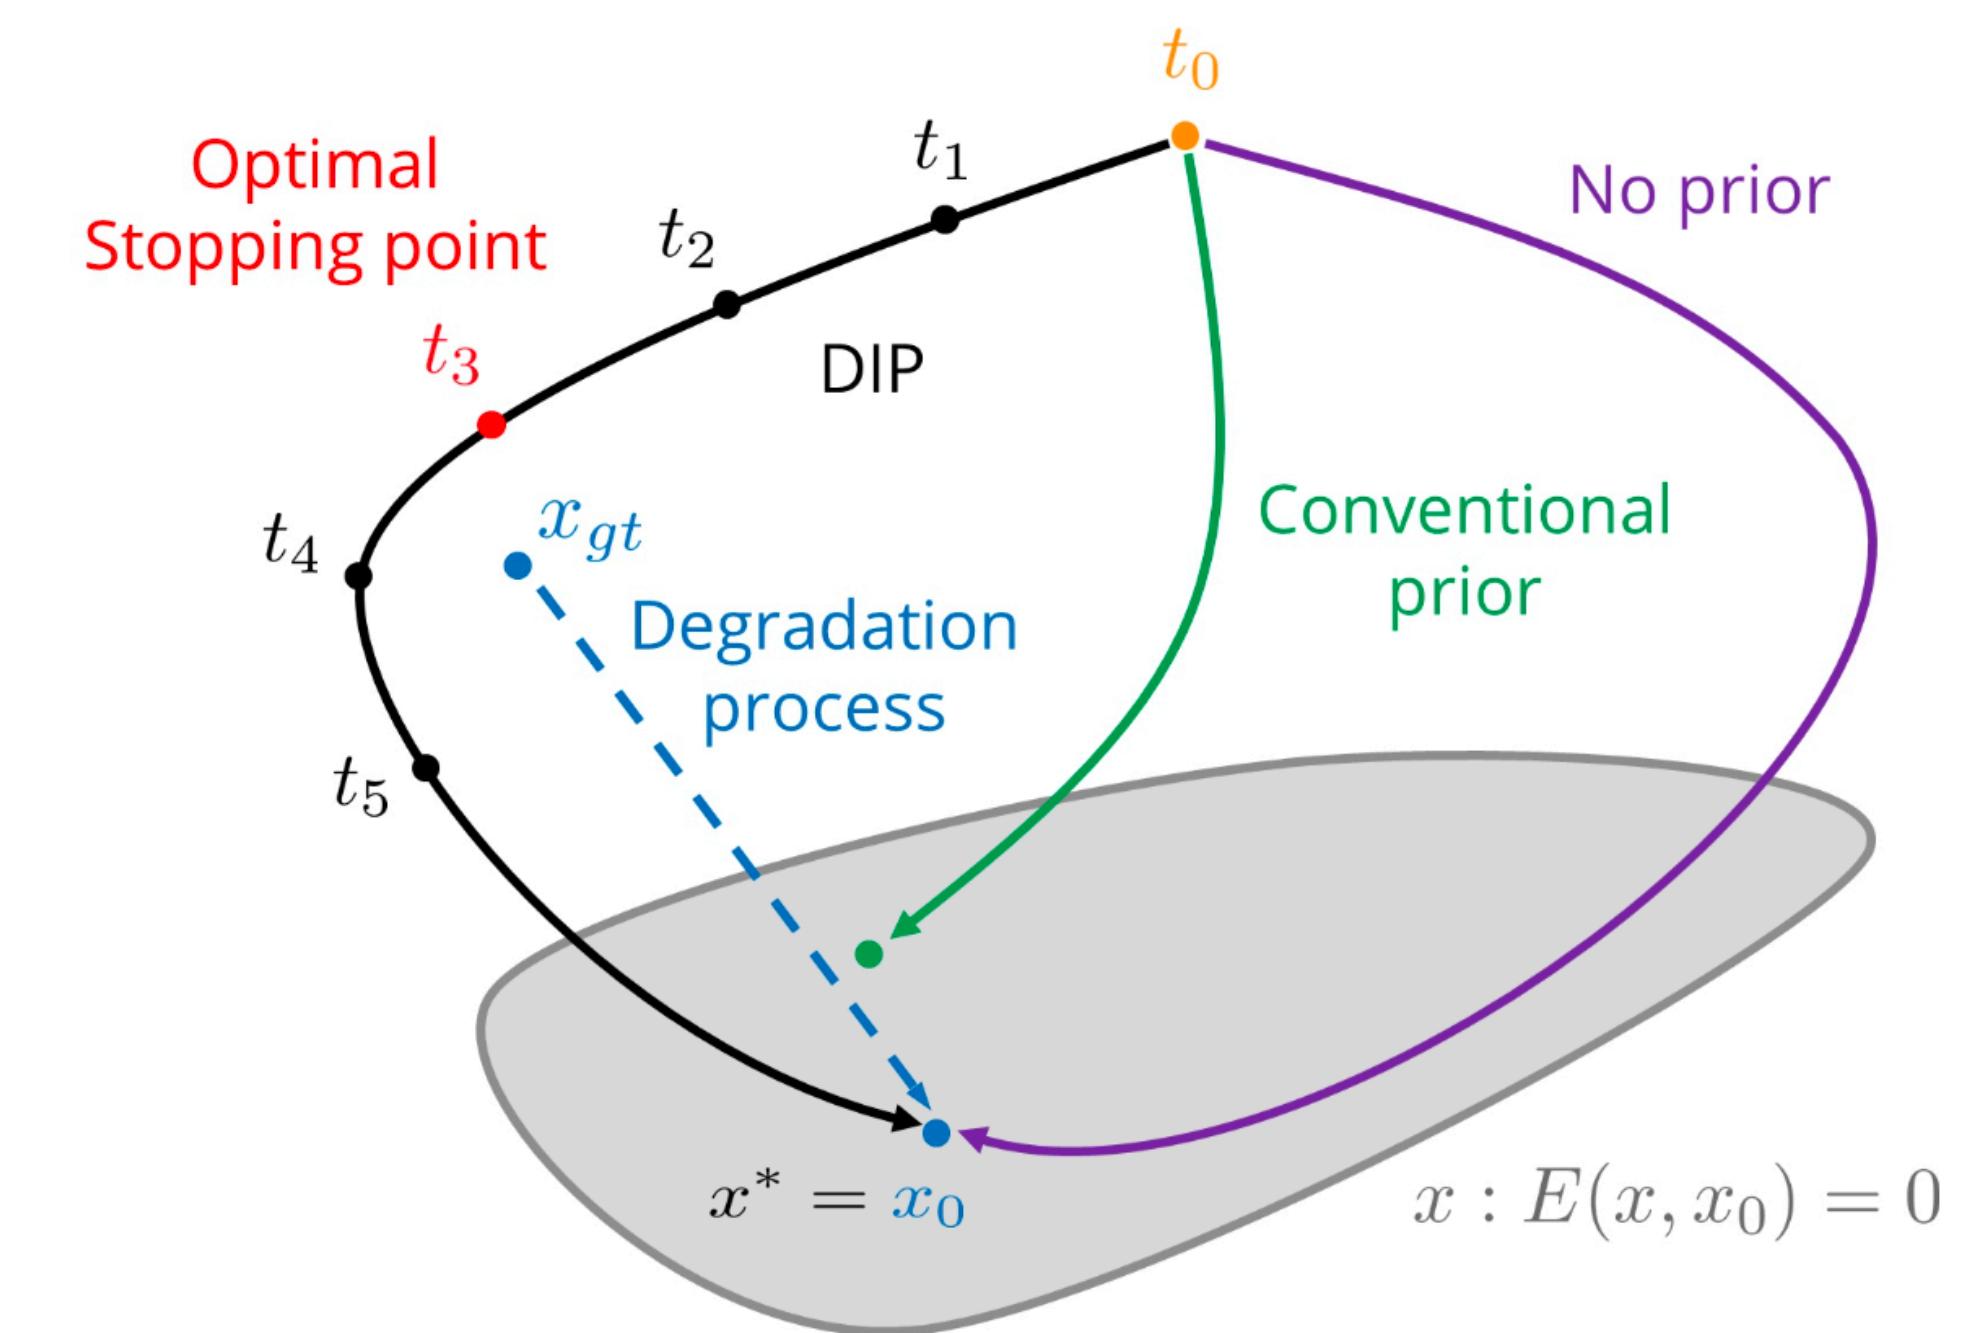
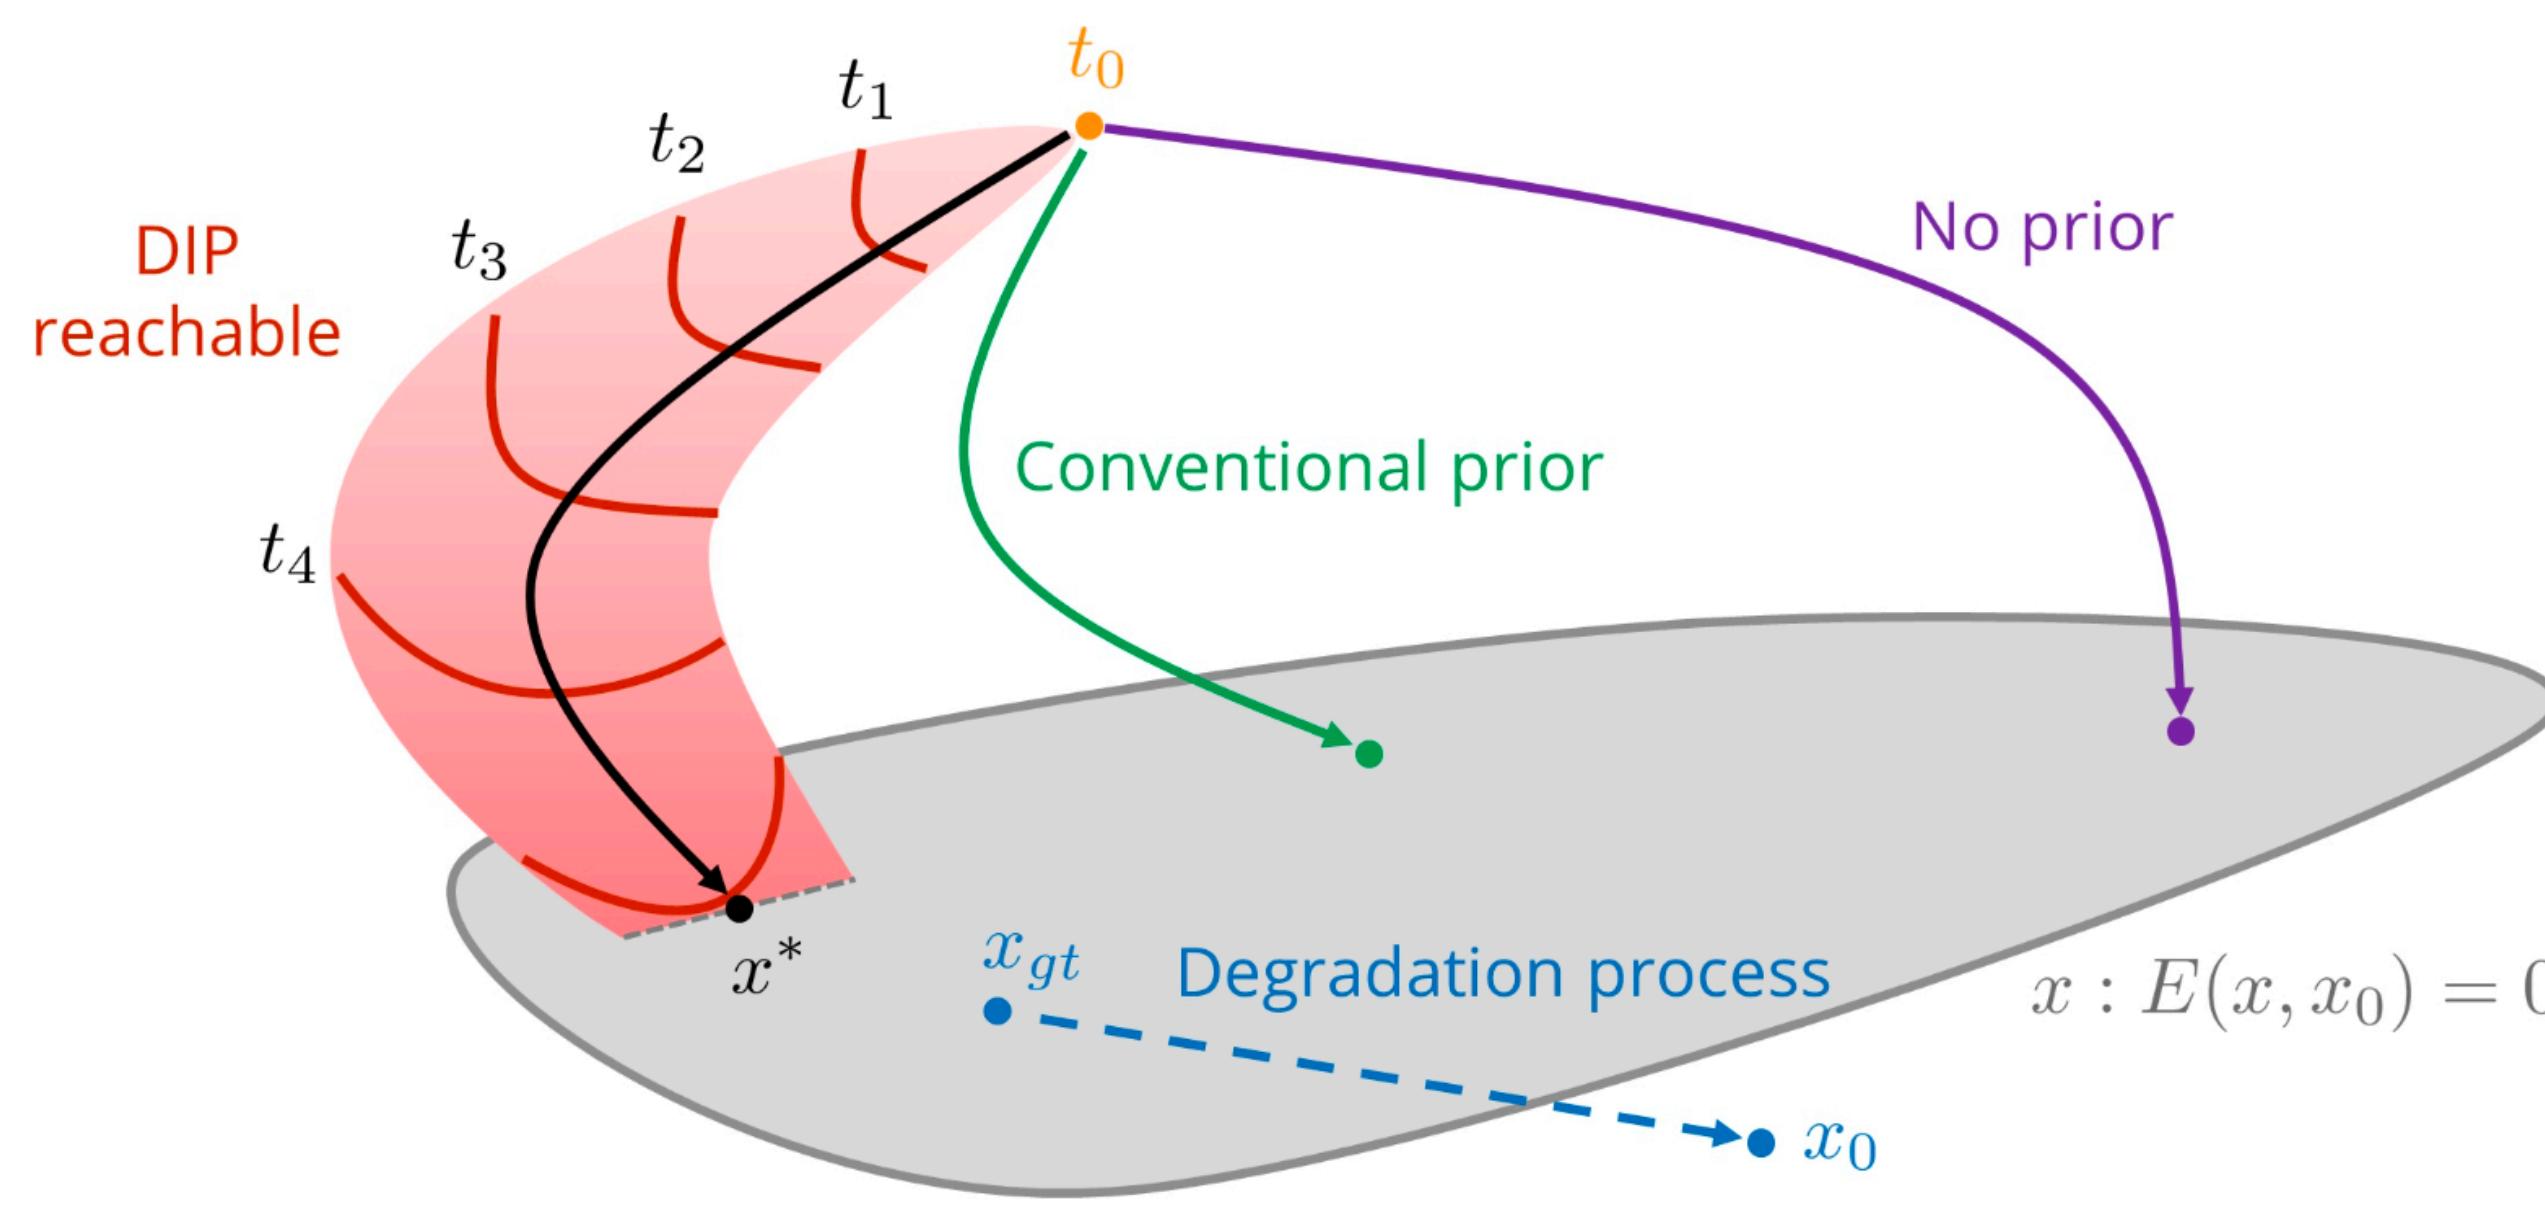
$$\theta^{k+1} = \theta^k - \alpha \frac{\partial E(f_{\theta}(z); \hat{x})}{\partial \theta}$$

3. Получаем решение $x^* = f_{\theta^*}(z)$

Алгоритм



Как работает prior в Deep Image Prior



Архитектура

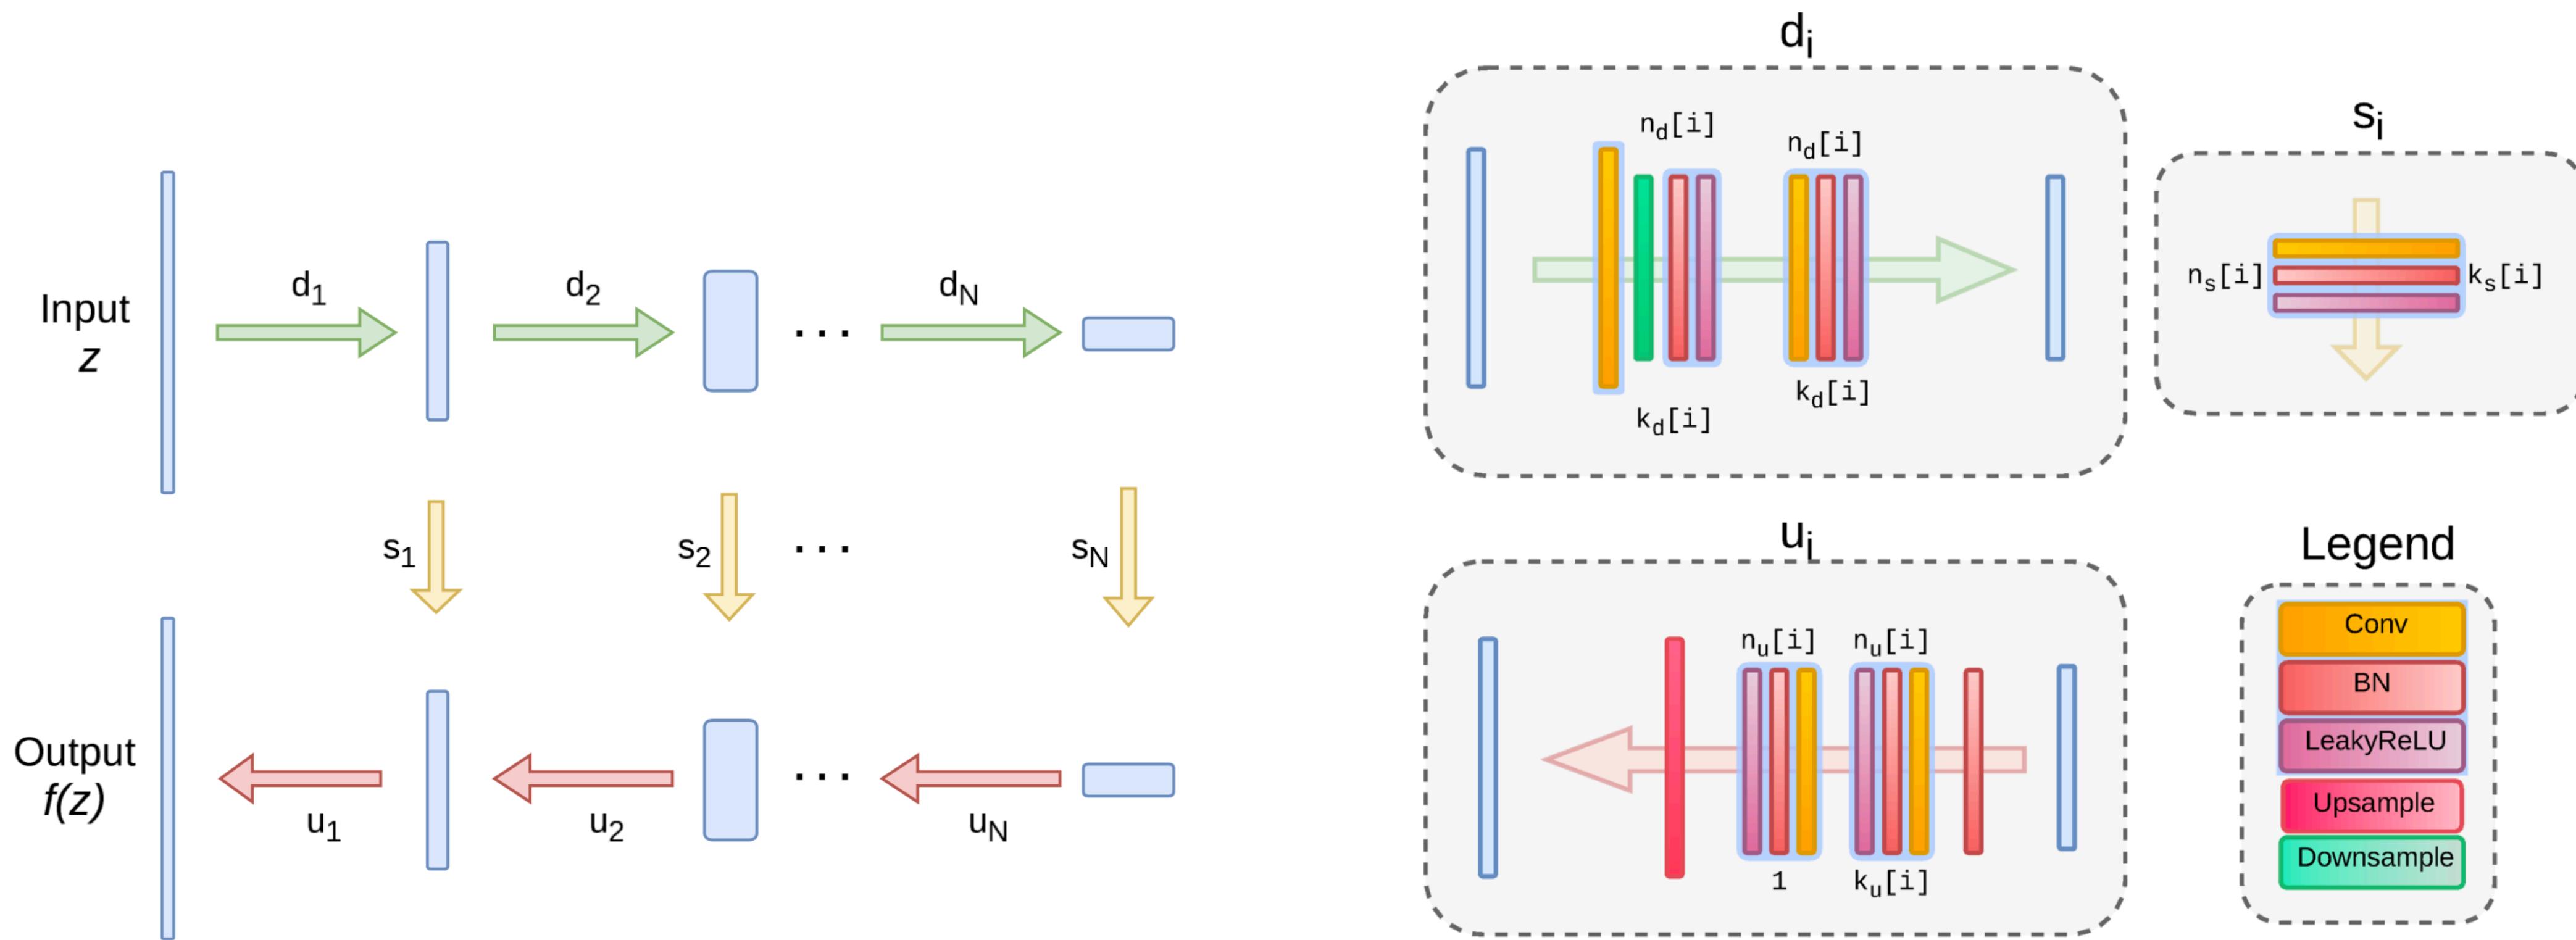


Fig. 21: **The architecture used in the experiments.** We use “hourglass” (also known as “decoder-encoder”) architecture. We sometimes add skip connections (yellow arrows). $n_u[i]$, $n_d[i]$, $n_s[i]$ correspond to the number of filters at depth i for the upsampling, downsampling and skip-connections respectively. The values $k_u[i]$, $k_d[i]$, $k_s[i]$ correspond to the respective kernel sizes.

Data terms

- **Denoising:** $E(x; \hat{x}) = \|x - \hat{x}\|^2$
- **Inpainting:** $E(x; \hat{x}) = \|(x - \hat{x}) \cdot m\|^2$, m - маска
- **Super-resolution:** $E(x; \hat{x}) = \|d(x) - \hat{x}\|^2$

Denoising

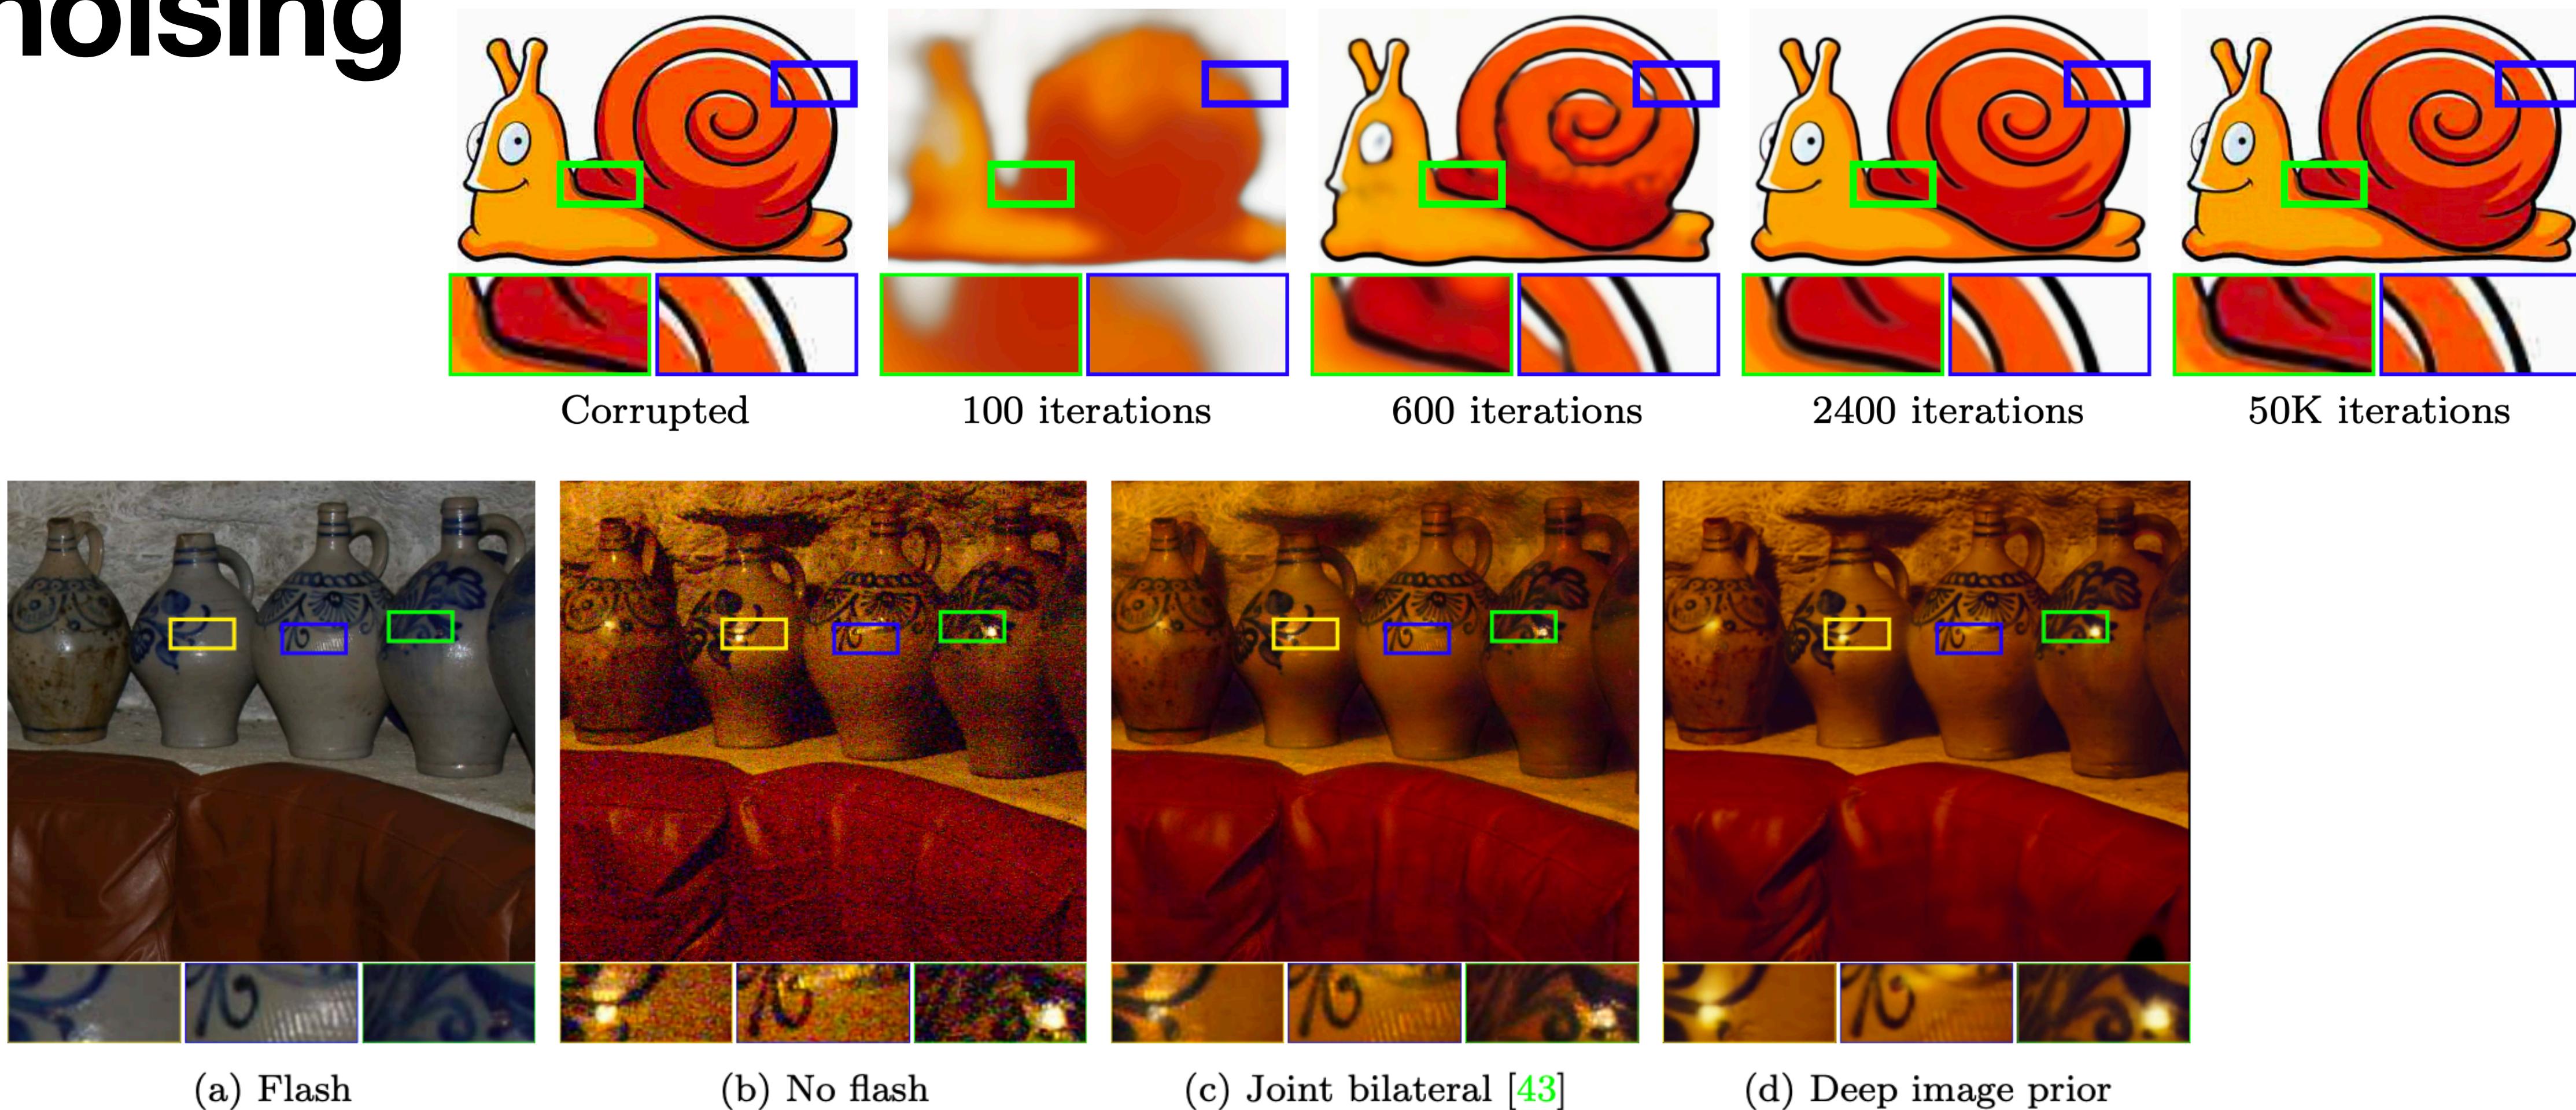
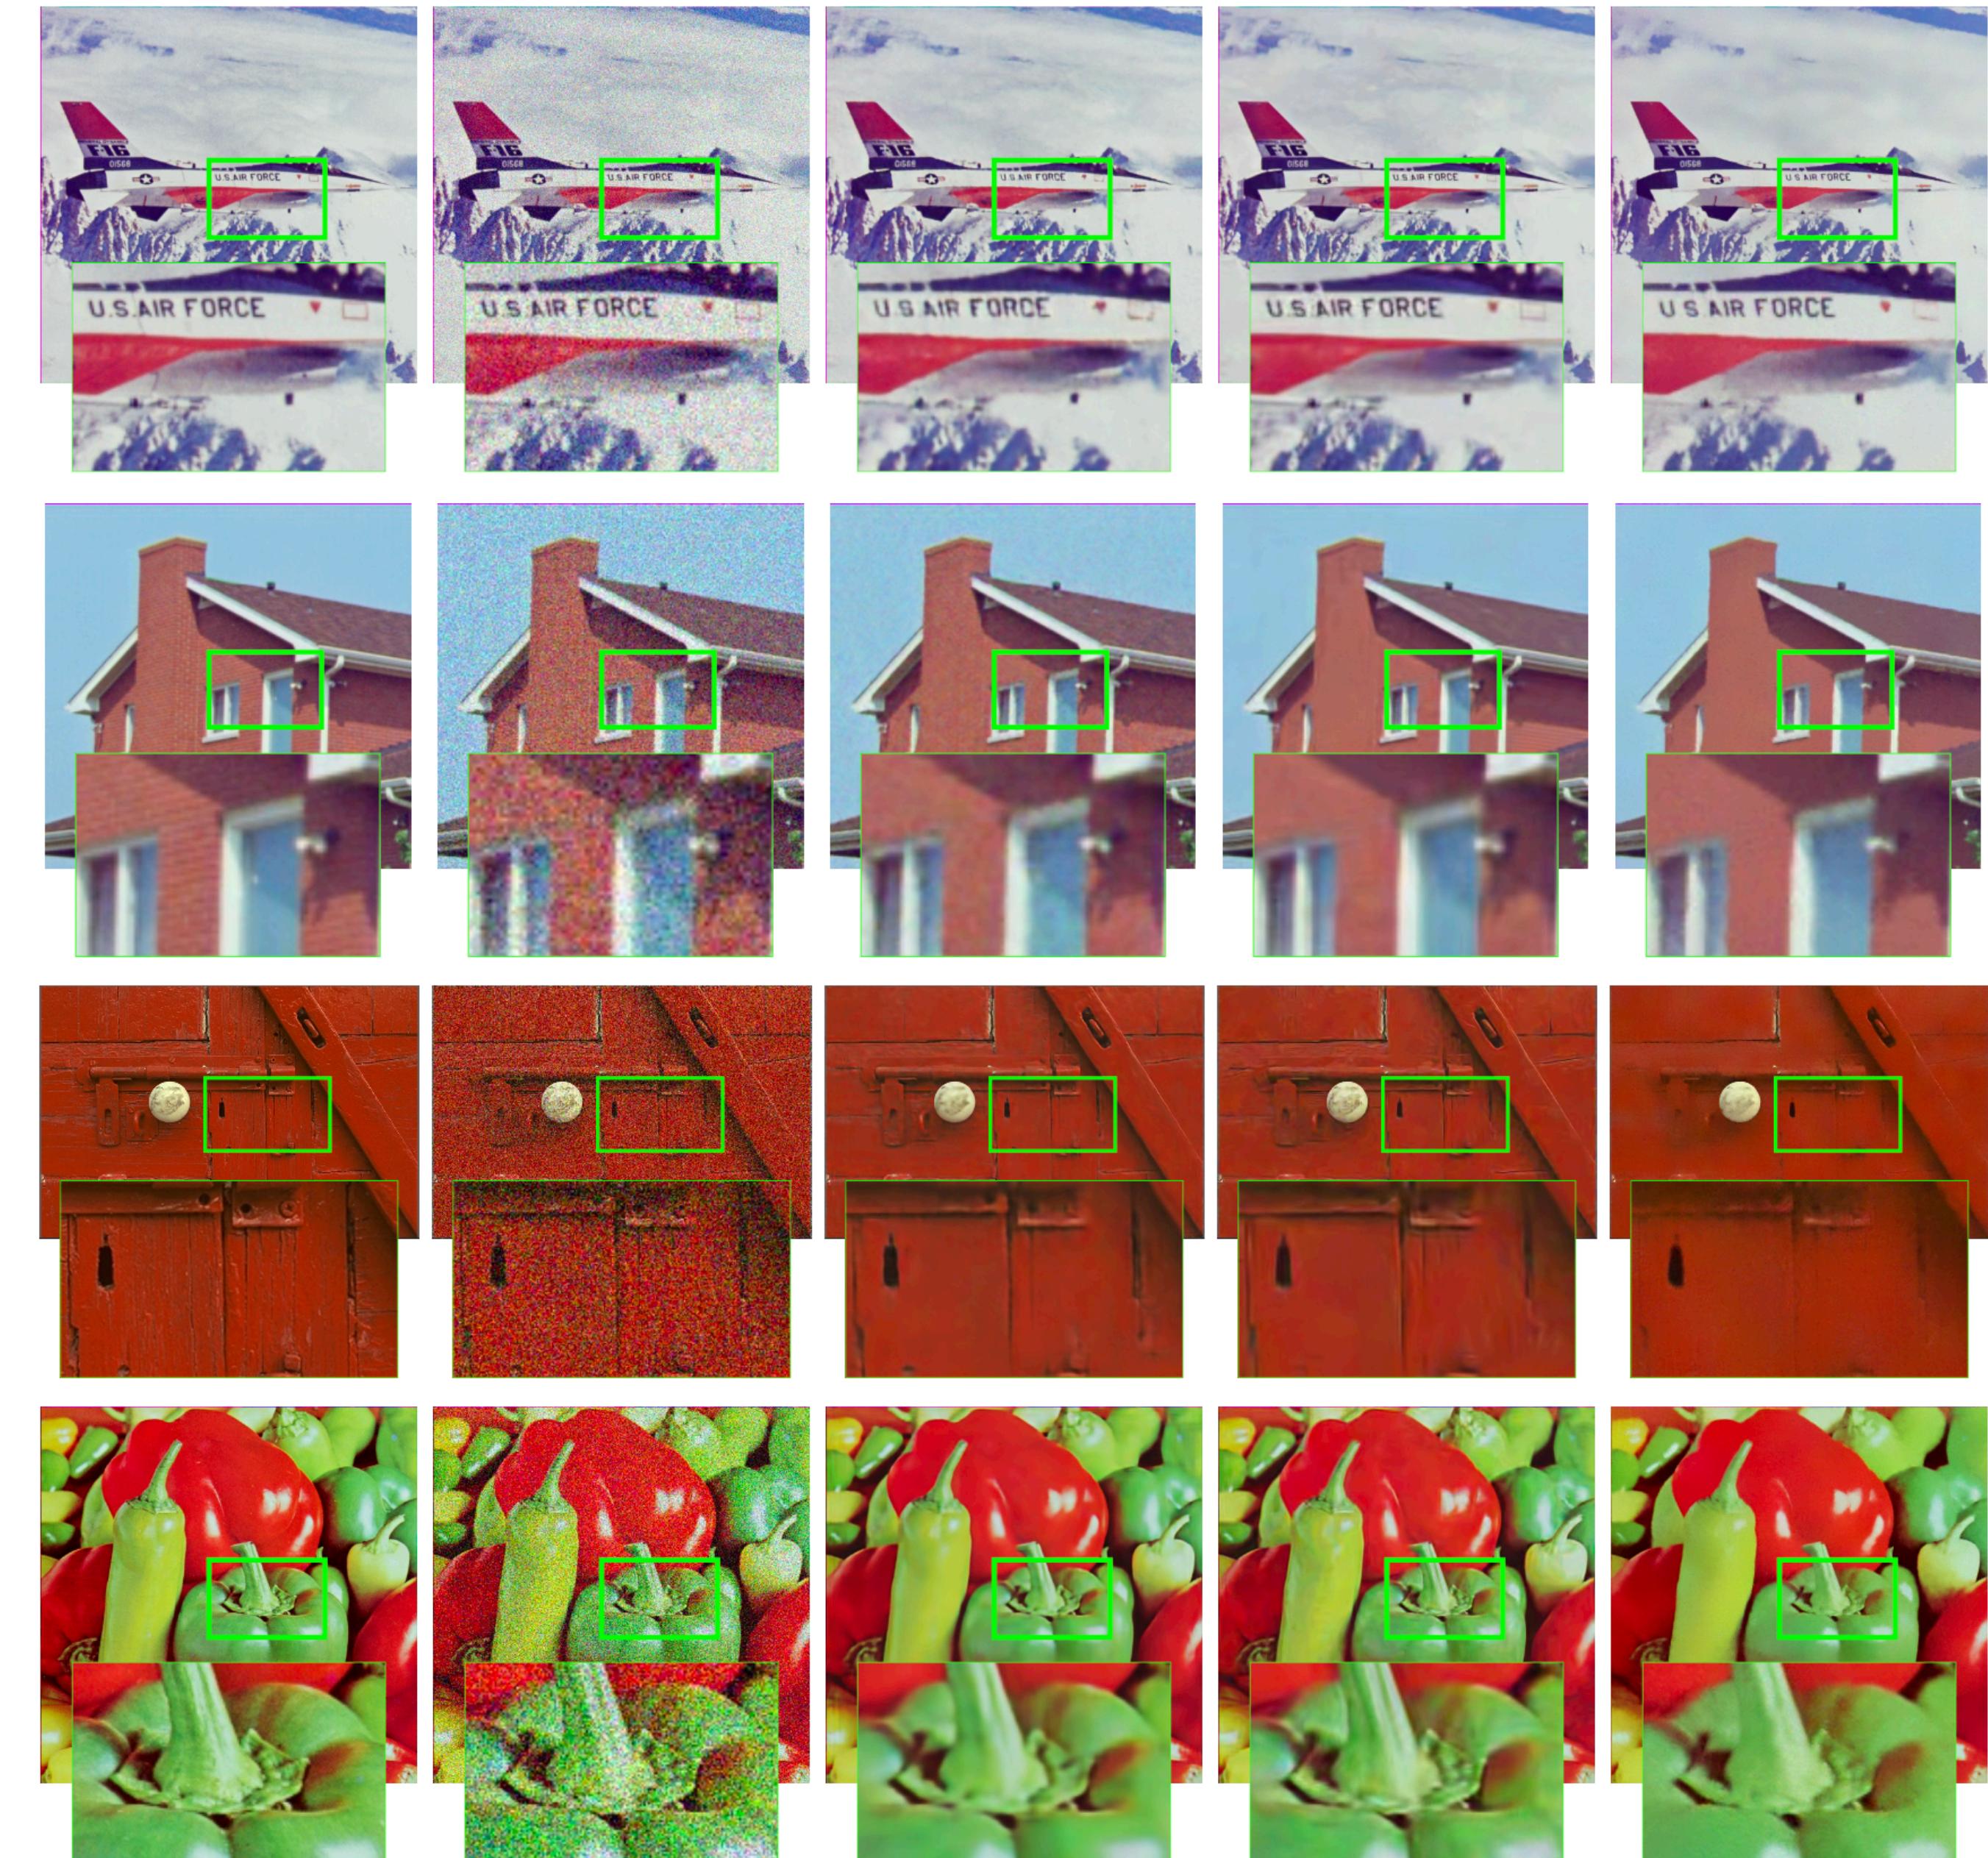


Fig. 20: **Reconstruction based on flash and no-flash image pair.** The deep image prior allows to obtain low-noise reconstruction with the lighting very close to the no-flash image. It is more successful at avoiding “leaks” of the lighting patterns from the flash pair than joint bilateral filtering [43] (c.f. blue inset).

Denoising



(a) GT

(b) Input

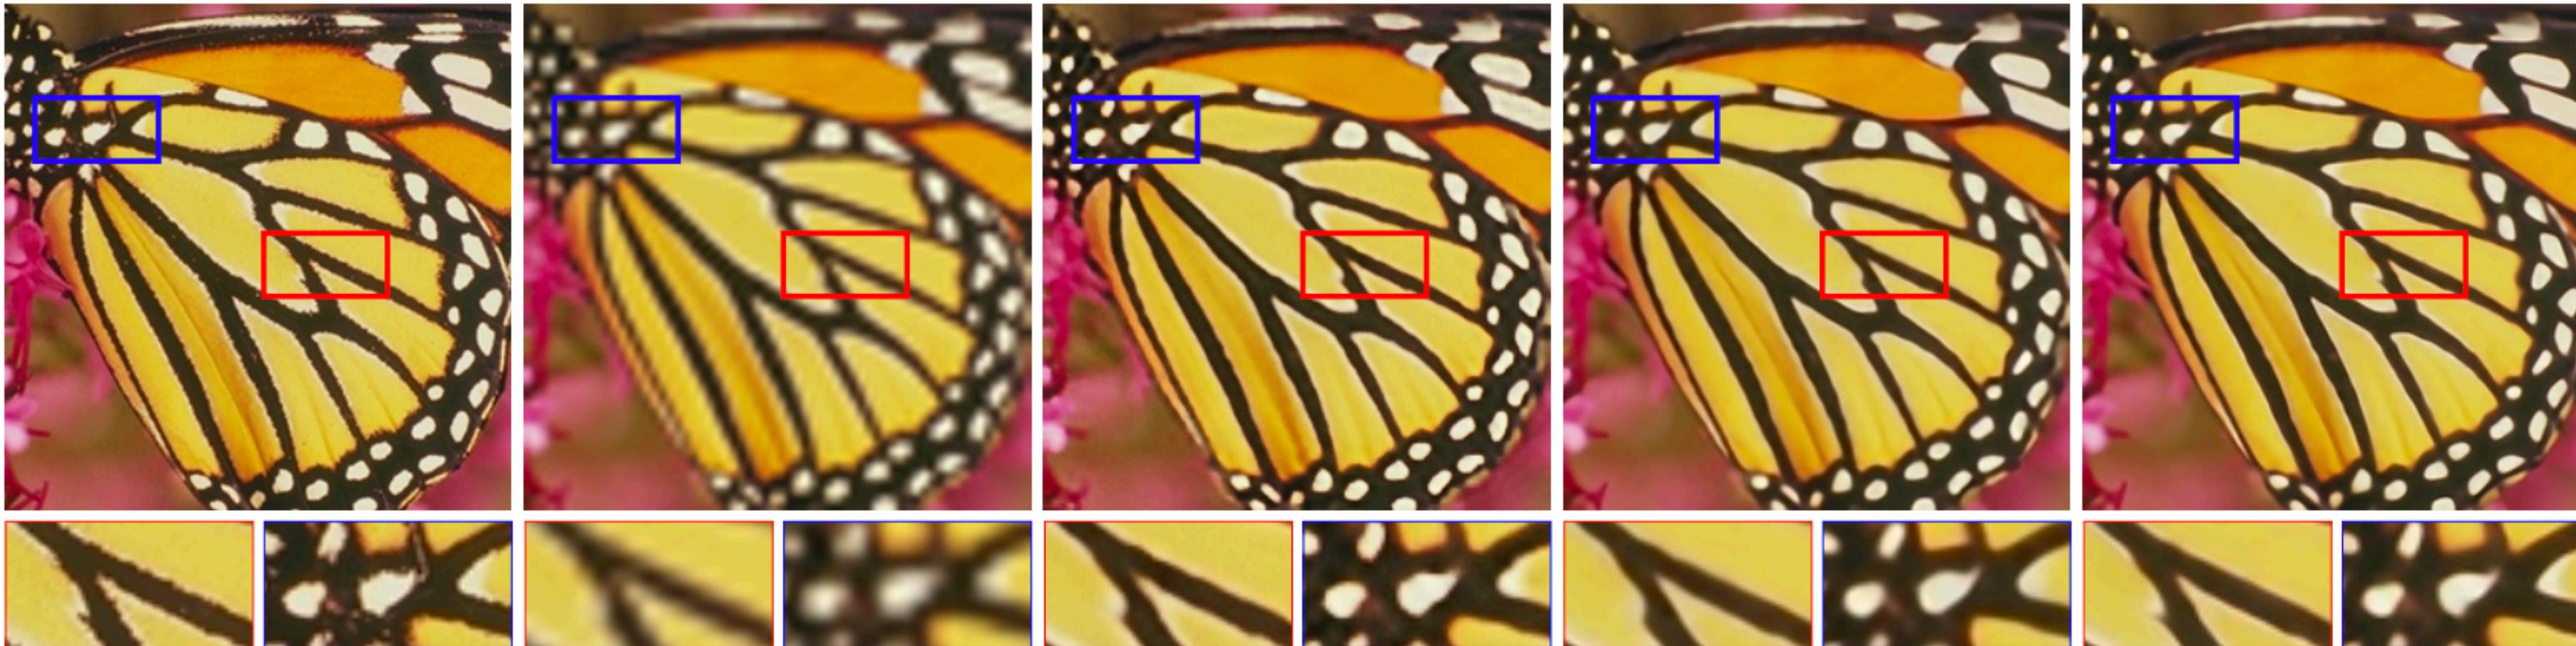
(c) Ours

(d) CBM3D

(e) NLM

Super-resolution

4× super-resolution



(a) Original image

(b) Bicubic,
Not trained

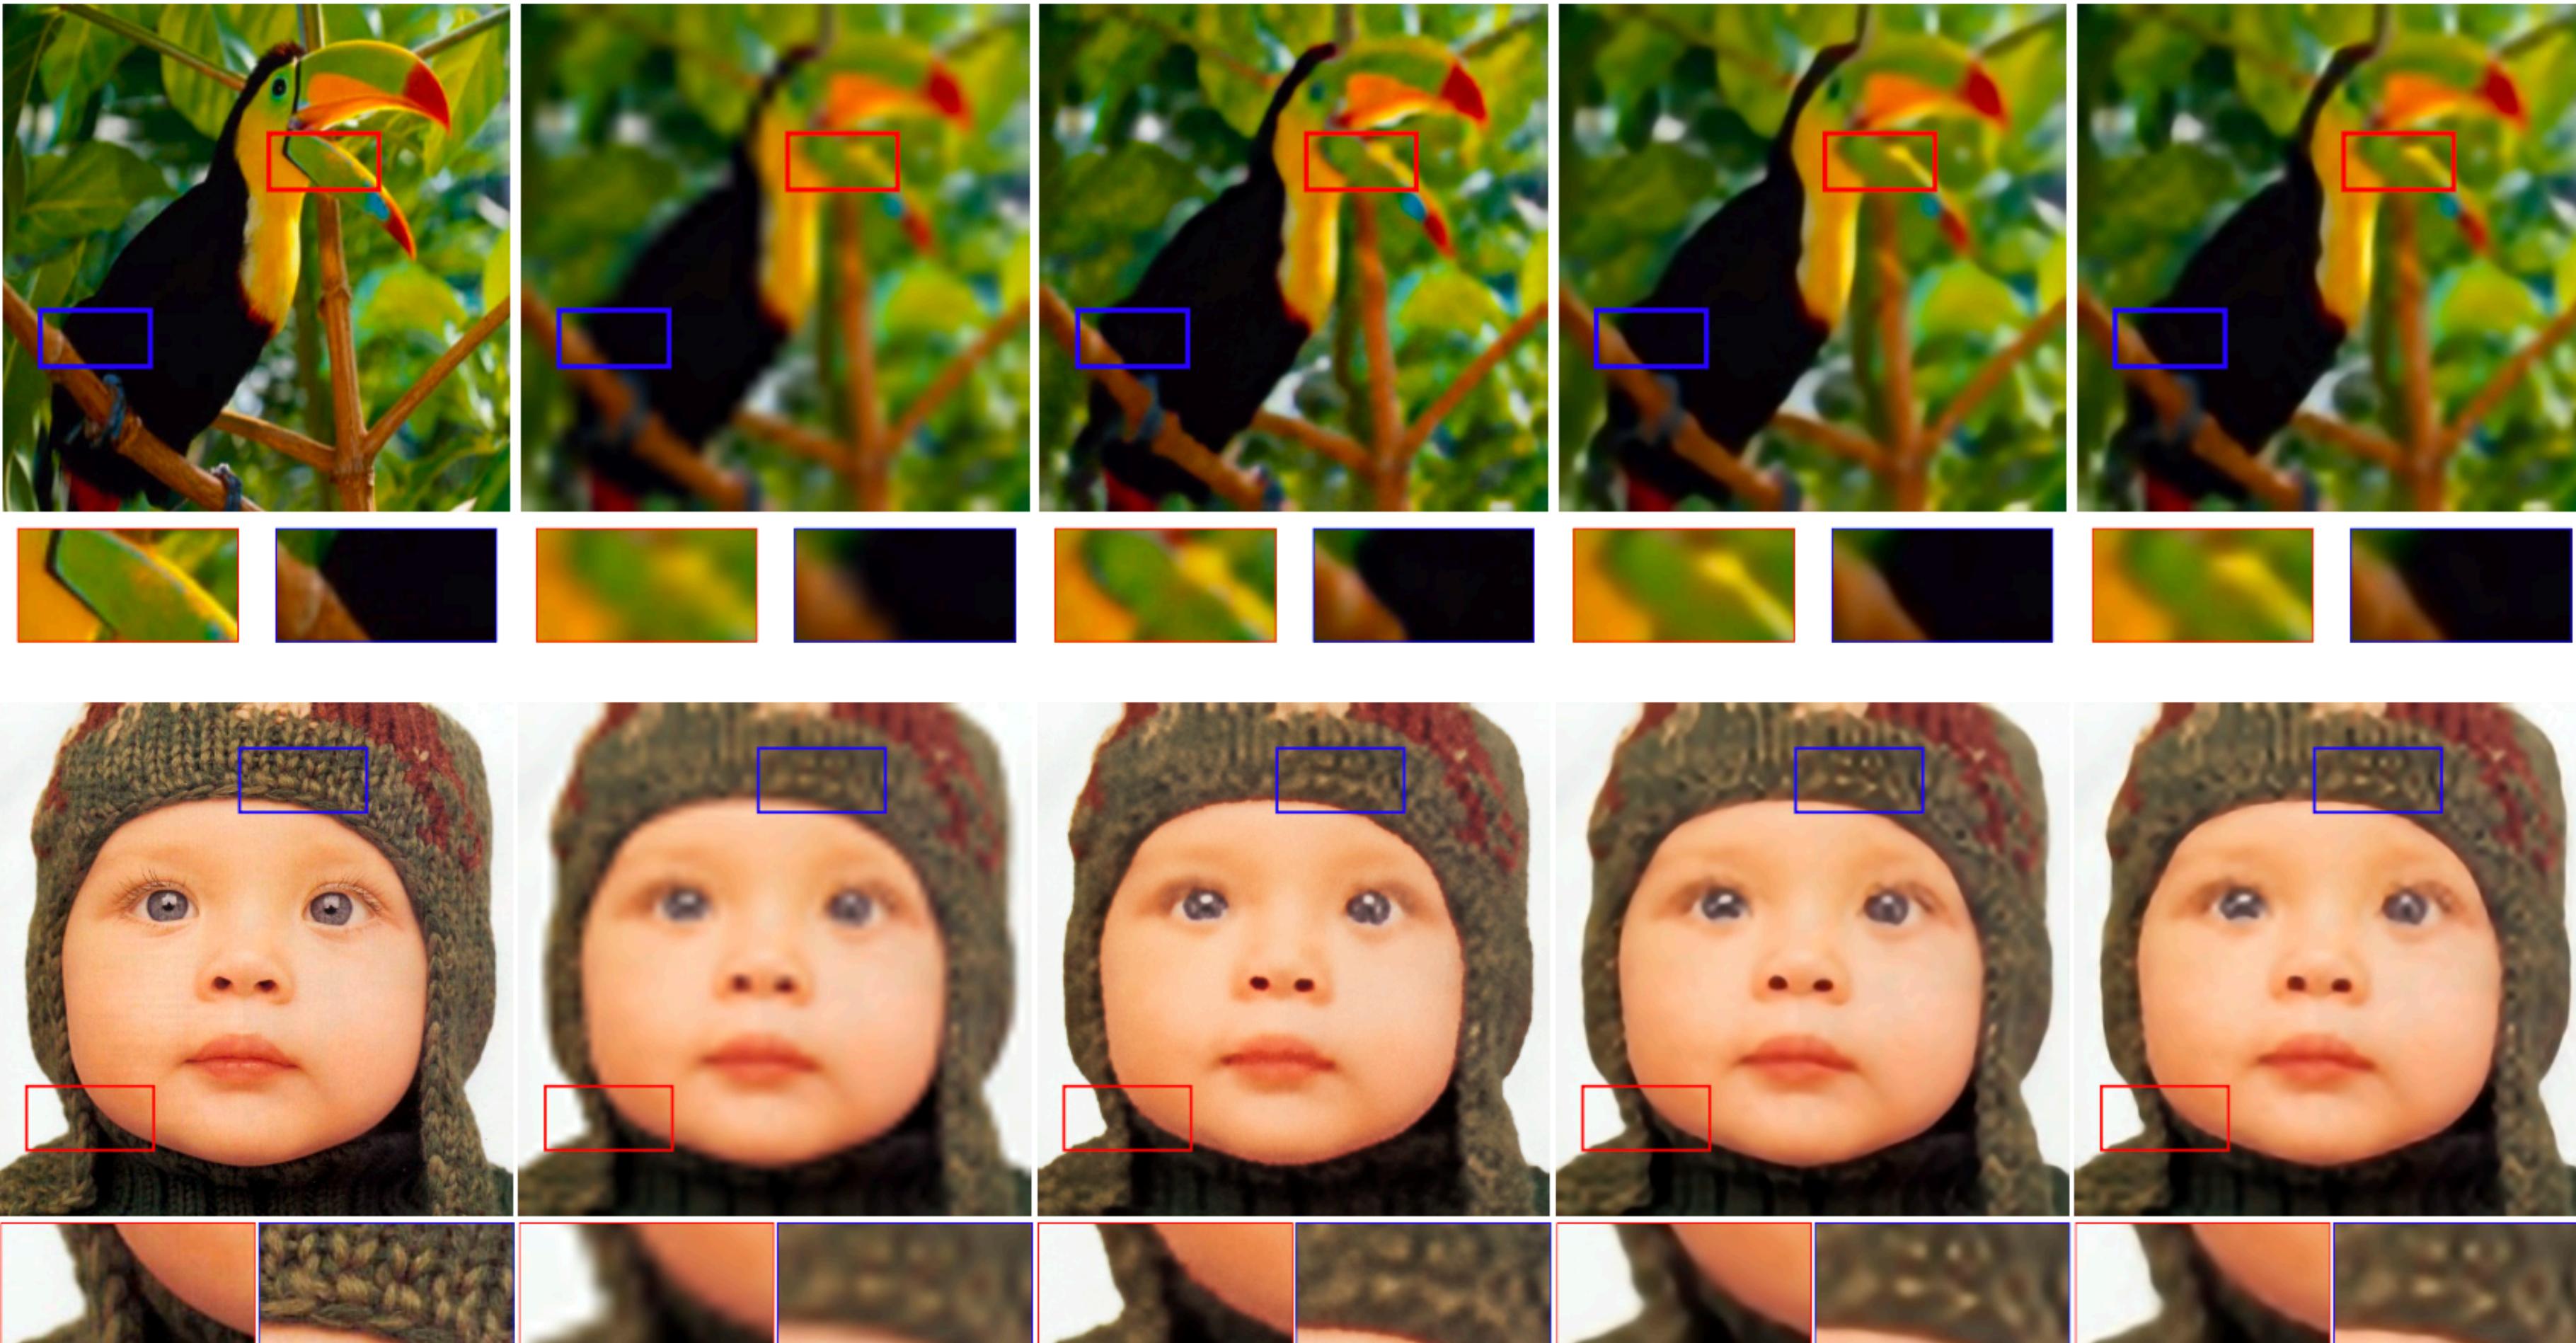
(c) Ours,
Not trained

(d) LapSRN,
Trained

(e) SRResNet,
Trained

Super-resolution

8× super-resolution



(f) Original image

(g) Bicubic,
Not trained

(h) Ours,
Not trained

(i) LapSRN,
Trained

(j) VDSR,
Trained

Inpainting



(a) Input (white=masked)



(b) Encoder-decoder, depth=6



(c) Encoder-decoder, depth=4



(d) Encoder-decoder, depth=2



(e) ResNet, depth=8



(f) U-net, depth=5

Inpainting

	Barbara	Boat	House	Lena	Peppers	C.man	Couple	Finger	Hill	Man	Montage
Papyan et al.	28.14	31.44	34.58	35.04	31.11	27.90	31.18	31.34	32.35	31.92	28.05
Ours	32.22	33.06	39.16	36.16	33.05	29.8	32.52	32.84	32.77	32.20	34.54

Fig. 12: Comparison between our method and the algorithm in [42]. See fig. 13 for visual comparison.

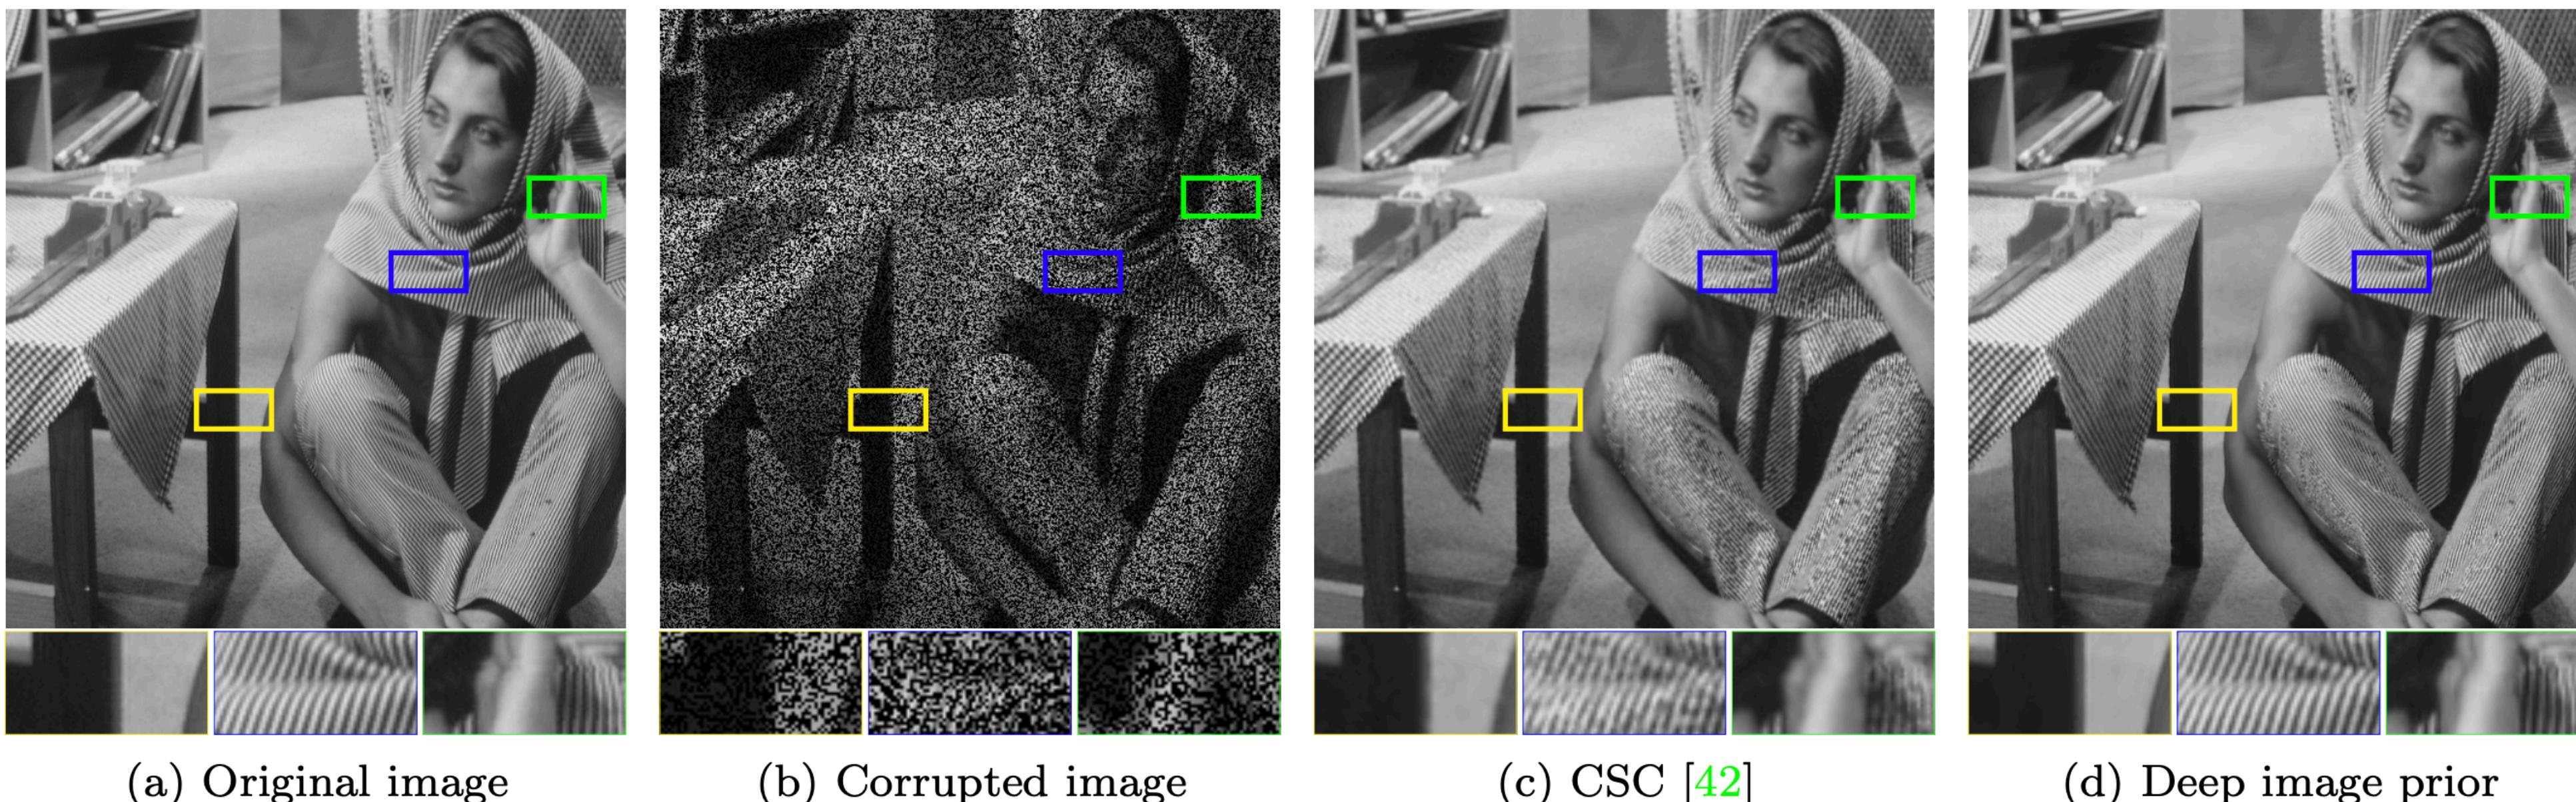
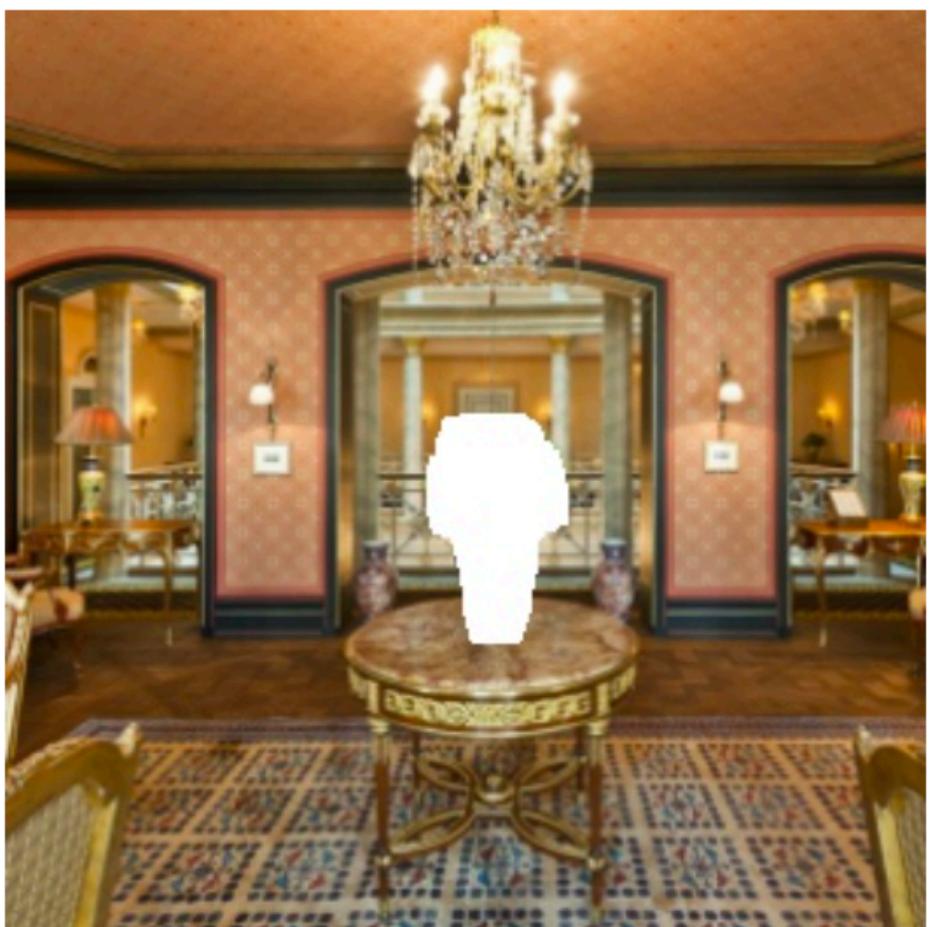


Fig. 13: Comparison with convolutional sparse coding (CSC) [42] on inpainting 50% of missing pixels. Our approach recovers a natural image with more accurate fine details than convolutional sparse coding.

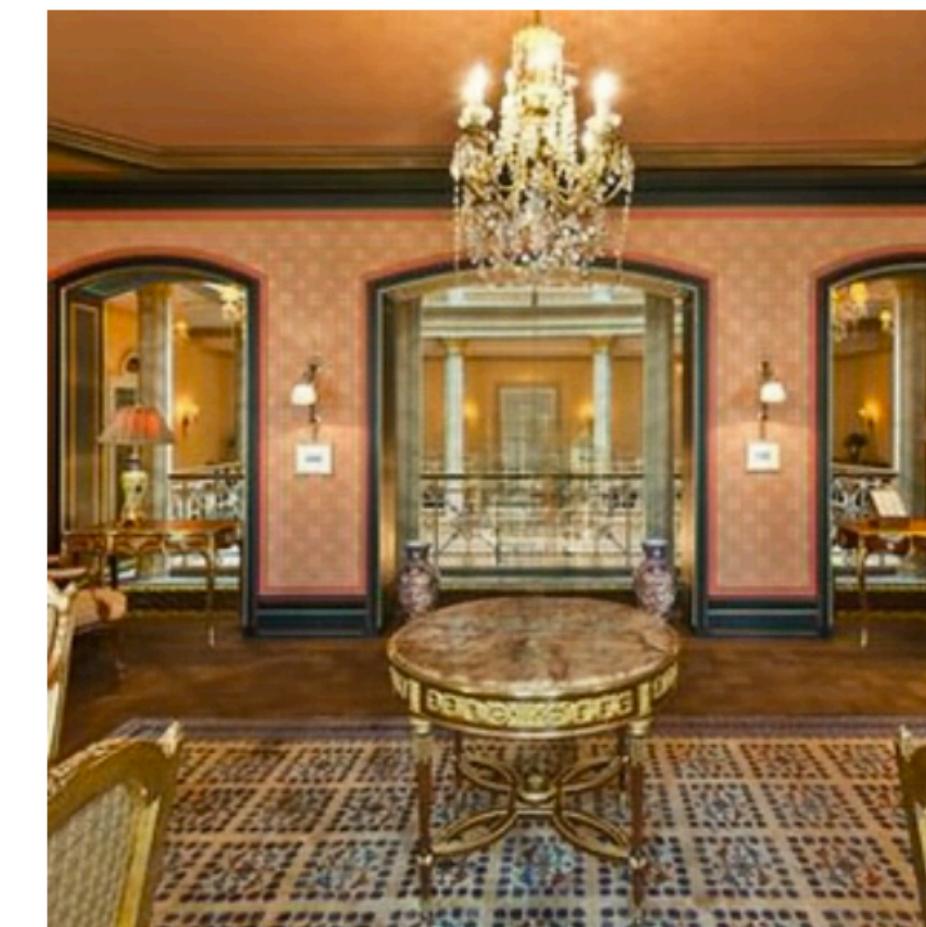
Inpainting



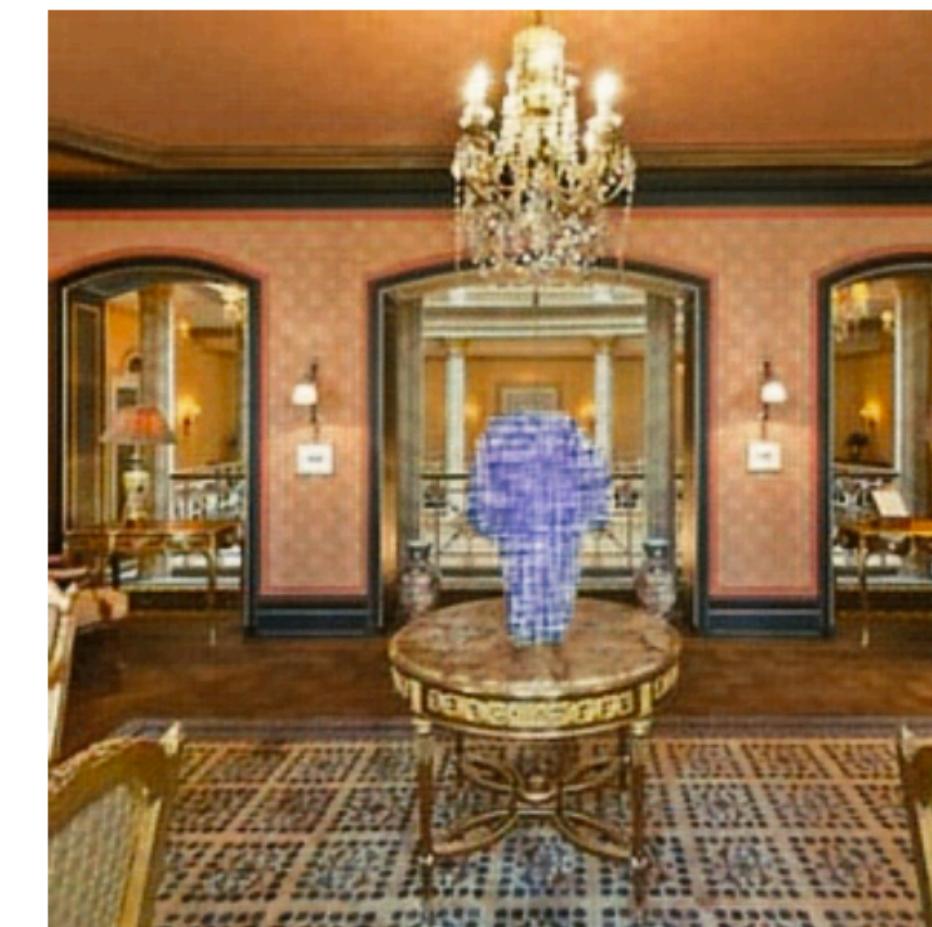
(a) Corrupted image



(b) Global-Local GAN [27]

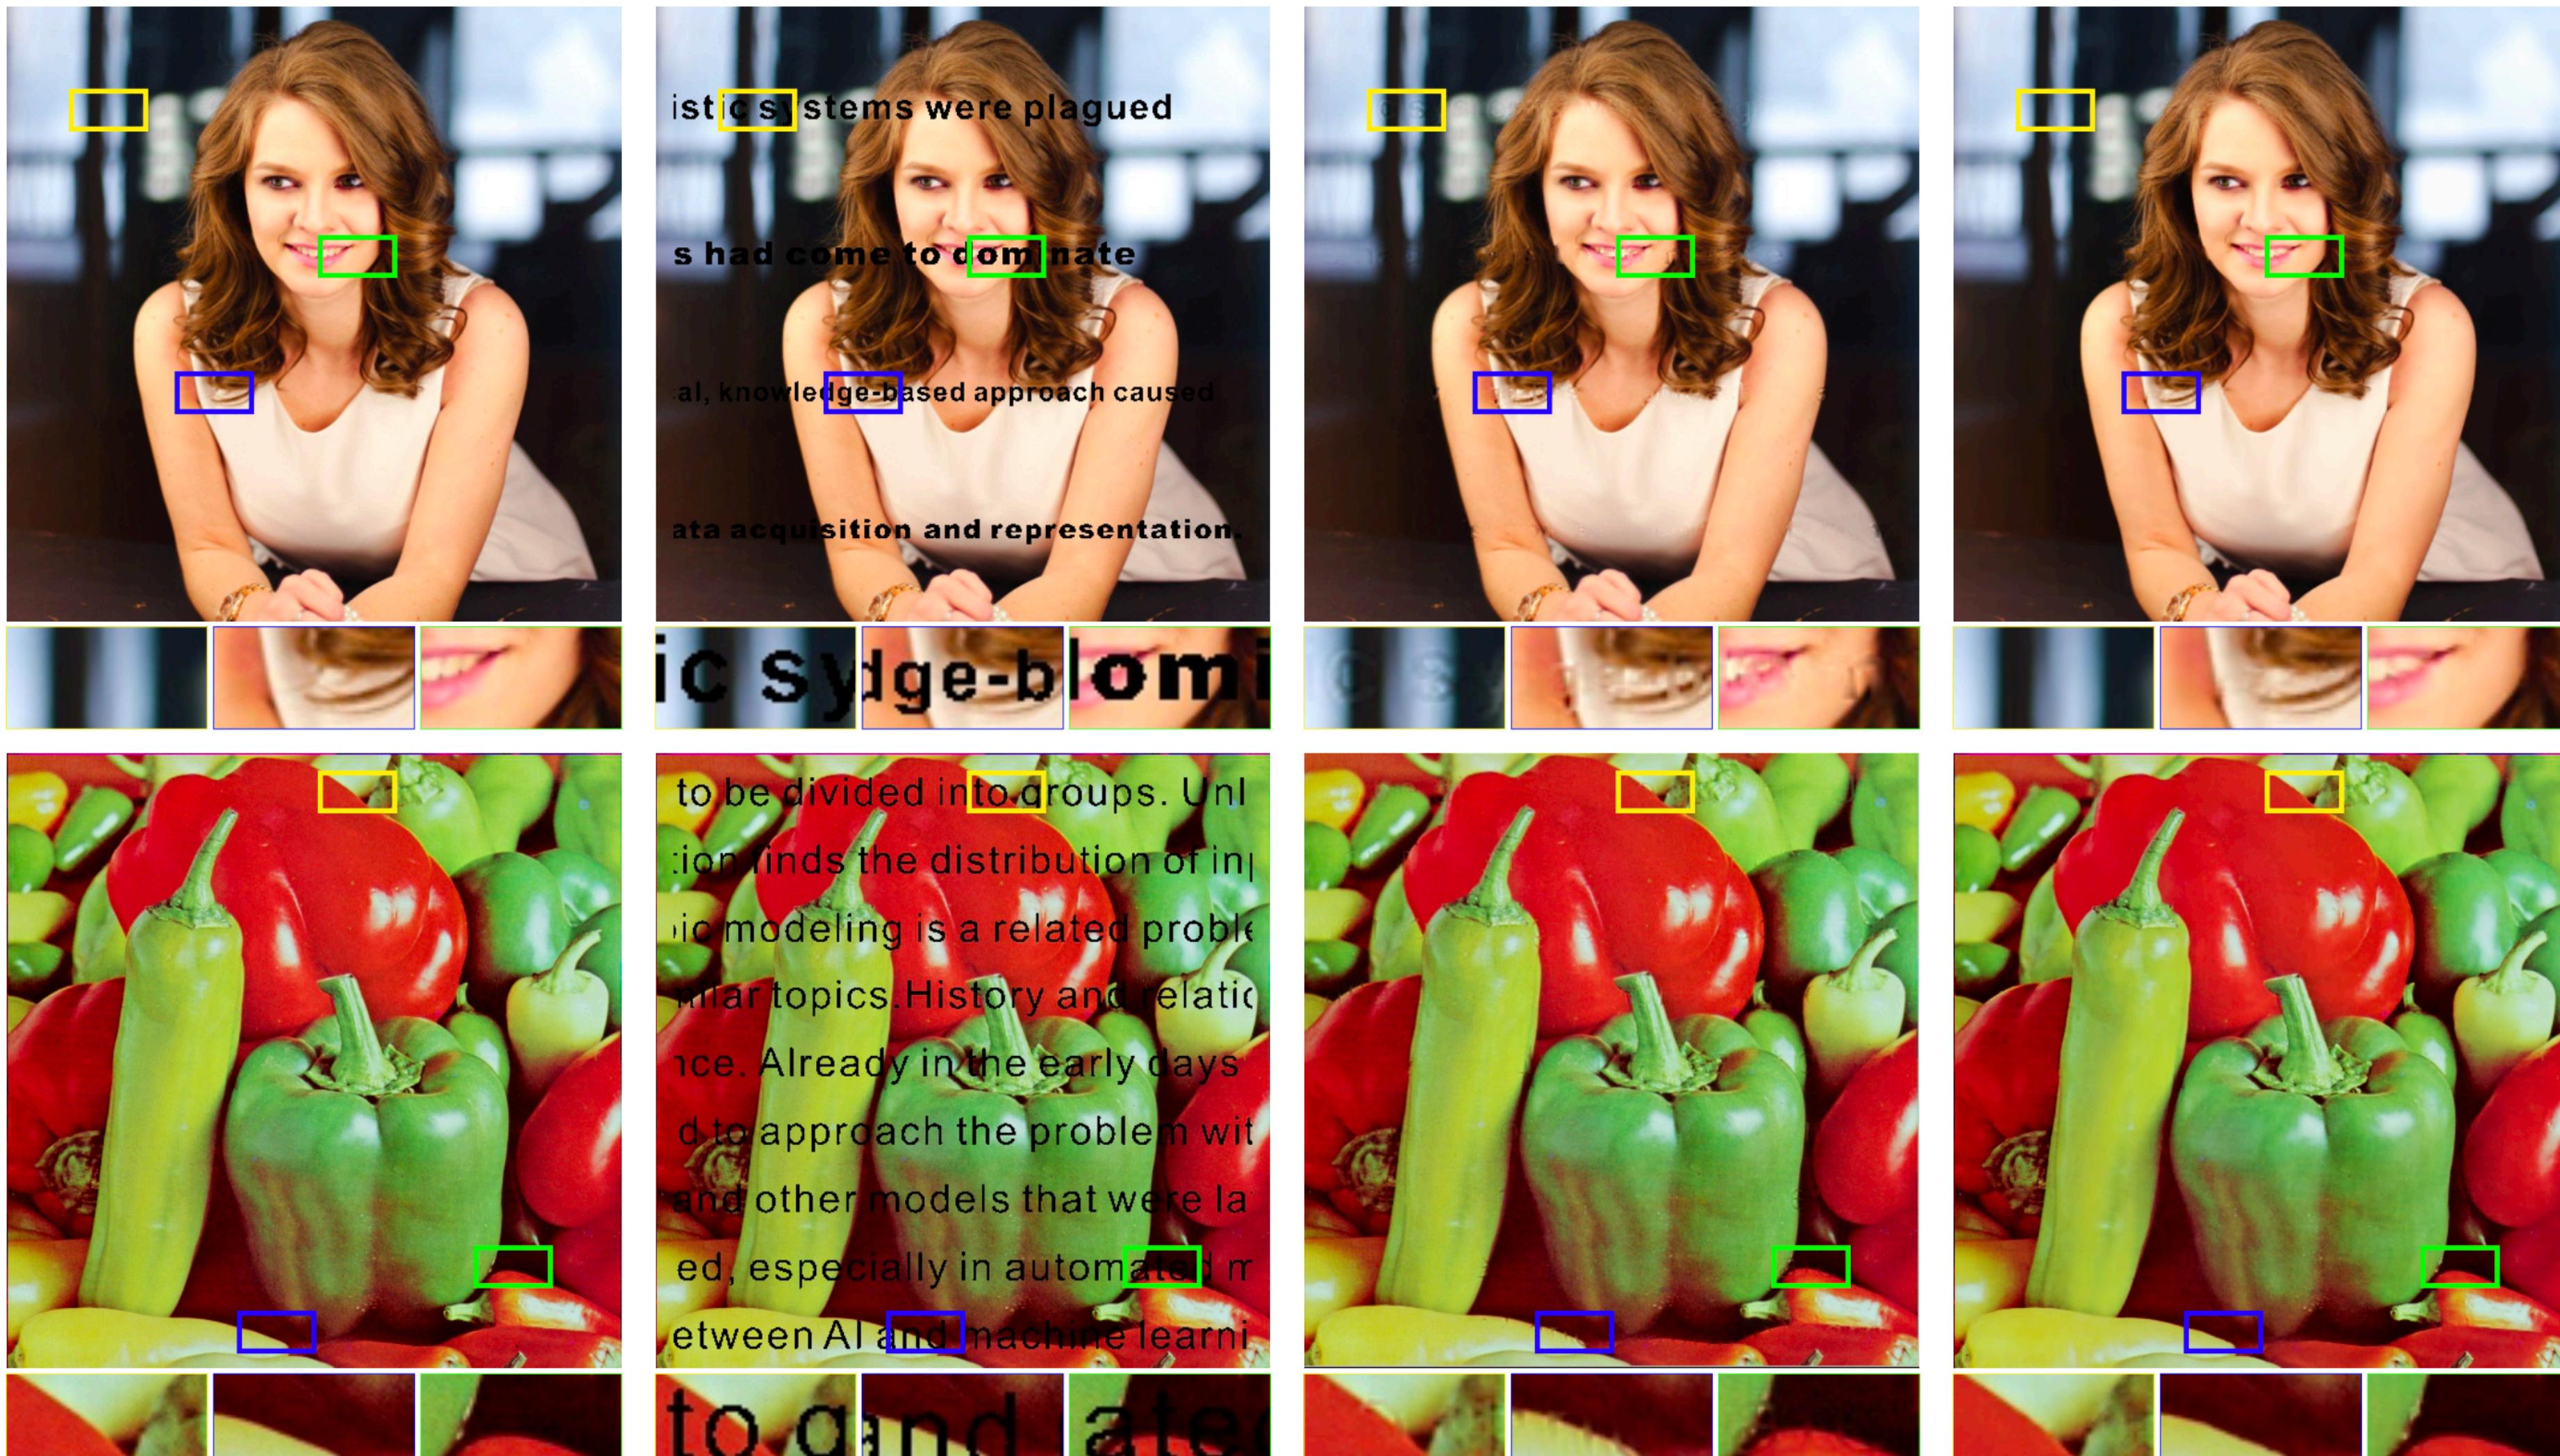


(c) Ours, $LR = 0.01$



(d) Ours, $LR = 10^{-4}$

Inpainting



(a) Original image

(b) Corrupted image

(c) Shepard networks [44]

(d) Deep Image Prior

Выводы

- Это не новый метод решения задачи image restoration.
- Получается, что архитектура нейросети сама по себе формирует prior, который обычно получается благодаря обучению.
- Работает долго и нужна специальное отслеживание количества итераций.
- Результаты работы нейросети без обучения близок к нейросетям, которые были обучены на датасете.