



# GAN - Generative Adversarial Network

Cumulative number of named GAN papers by month






Labels to Street Scene



Labels to Facade



BW to Color



Aerial to Map



Day to Night





Edges to Photo






<https://arxiv.org/pdf/1611.07004>




$N(0, 1)$



$P_{fake}$



fake

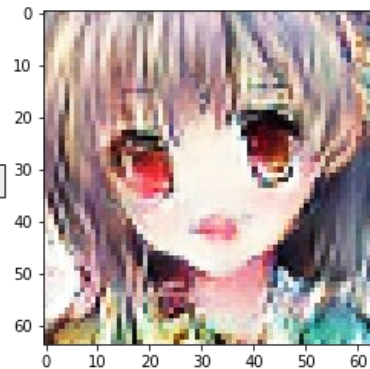


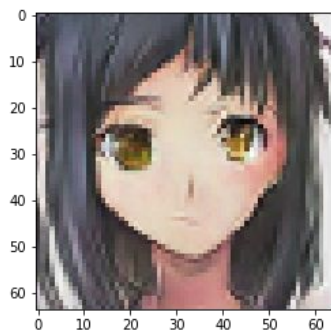









$P_{real}$




real





# DCGAN



# Checkerboard artefacts








Table 9: BigGAN-deep architecture for  $512 \times 512$  images.



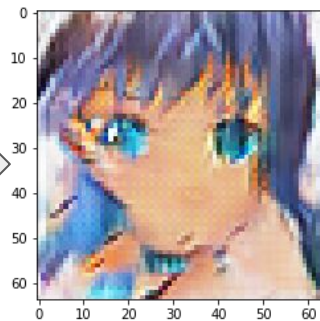
$z \in \mathbb{R}^{128} \sim \mathcal{N}(0, I)$ $\text{Embed}(y) \in \mathbb{R}^{128}$	$\text{RGB image } x \in \mathbb{R}^{512 \times 512 \times 3}$
Linear $(128 + 128) \rightarrow 4 \times 4 \times 16ch$	$3 \times 3 \text{ Conv } 3 \rightarrow ch$
ResBlock $16ch \rightarrow 16ch$	ResBlock down $ch \rightarrow ch$
ResBlock up $16ch \rightarrow 16ch$	ResBlock $ch \rightarrow ch$
ResBlock $16ch \rightarrow 16ch$	ResBlock down $ch \rightarrow 2ch$
ResBlock up $16ch \rightarrow 8ch$	ResBlock $2ch \rightarrow 2ch$
ResBlock $8ch \rightarrow 8ch$	ResBlock down $2ch \rightarrow 4ch$
ResBlock up $8ch \rightarrow 8ch$	ResBlock $4ch \rightarrow 4ch$
ResBlock $8ch \rightarrow 8ch$	Non-Local Block $(64 \times 64)$
ResBlock up $8ch \rightarrow 4ch$	ResBlock down $4ch \rightarrow 8ch$
Non-Local Block $(64 \times 64)$	ResBlock $8ch \rightarrow 8ch$
ResBlock $4ch \rightarrow 4ch$	ResBlock down $8ch \rightarrow 8ch$
ResBlock up $4ch \rightarrow 2ch$	ResBlock $8ch \rightarrow 8ch$
ResBlock $2ch \rightarrow 2ch$	ResBlock down $8ch \rightarrow 16ch$
ResBlock up $2ch \rightarrow ch$	ResBlock $16ch \rightarrow 16ch$
ResBlock $ch \rightarrow ch$	ResBlock down $16ch \rightarrow 16ch$
ResBlock up $ch \rightarrow ch$	ResBlock $16ch \rightarrow 16ch$
BN, ReLU, $3 \times 3 \text{ Conv } ch \rightarrow 3$	ReLU, Global sum pooling
Tanh	$\text{Embed}(y) \cdot h + (\text{linear} \rightarrow 1)$
(a) Generator	(b) Discriminator

1. Нормализация картинки тангенсом
2.  $\min \log(1-D) \rightarrow -\log(D)$
3. Нормальное распределение
4. BatchNorm
5. Без нулевых градиентов (ReLU, MaxPool)
6. Noisy labels
7. DCGAN
8. Параметры Adam


$N(0, 1)$



$P_{fake}$



fake





$P_{real}$






real

$N(0, 1)$




$P_{fake}$




fake

real



$P_{real}$



$$Loss = E_{x \in P_{real}} [\log D(x)] + E_{x \in P_{fake}} [\log(1 - D(x))]$$

Дискриминатор =  $\operatorname{argmax} (Loss)$

Генератор =  $\operatorname{argmin} (Loss)$

Пусть задан генератор G, каков дискриминатор D?

$$Loss = \int_x P_{real}(x) [\log D(x)] + \int_x P_{fake}(x) [\log(1 - D(x))] \rightarrow \max$$

$$Loss = \int_x P_{real}(x) [\log D(x)] + P_{fake}(x) [\log(1 - D(x))] \rightarrow \max$$

Интеграл - сумма по всем x, значит достаточно максимизировать каждый x в отдельности



Пусть задан генератор  $G$  и точка  $x$ , каков дискриминатор  $D(x)$ ?

$$P_{real}(x)[\log D(x)] + P_{fake}(x)[\log(1 - D(x))] \rightarrow \max$$

Берем производную по  $D$  приравниваем 0

$$P_{real}(x) \frac{1}{D(x)} + P_{fake}(x) \frac{-1}{1 - D(x)} = 0$$

$$(1 - D(x))P_{real}(x) = D(x)P_{fake}(x)$$

$$P_{real}(x) - D(x)P_{real}(x) = D(x)P_{fake}(x)$$

$$P_{real}(x) = D(x)(P_{fake}(x) + P_{real}(x))$$

$$D(x) = \frac{P_{real}(x)}{P_{fake}(x) + P_{real}(x)}$$

# Подставляем в исходное выражение

Получаем оптимальный дискриминатор

$$\int_x P_{real}(x) \left[ \log \frac{P_{real}(x)}{P_{fake}(x) + P_{real}(x)} \right] + P_{fake}(x) \left[ \log \frac{P_{fake}(x)}{P_{fake}(x) + P_{real}(x)} \right]$$




Осталось лишь привести его к красивому виду

$$\begin{aligned} -2 \log 2 + \int_x P_{real}(x) \left[ \log \frac{P_{real}(x)}{(P_{fake}(x) + P_{real}(x))/2} \right] + \int_x P_{fake}(x) \left[ \log \frac{P_{fake}(x)}{(P_{fake}(x) + P_{real}(x))/2} \right] \\ -2 \log 2 + KL \left( P_{real} \left| \frac{P_{real} + P_{fake}}{2} \right. \right) + KL \left( P_{fake} \left| \frac{P_{real} + P_{fake}}{2} \right. \right) \end{aligned}$$

$$D = -2 \log 2 + 2JS(P_{real}|P_{fake})$$

$$GAN = \operatorname{argmin}_G \operatorname{argmax}_D \operatorname{Loss}(P_{real}, P_{fake}) = \operatorname{argmin}_G JS(P_{real} | P_{fake})$$





# Проблема JS divergence






Earth mover distance






$$EMD(P_r, P_\theta) = \min_S \sum_{x,y} |x - y| S(x, y)$$





$\Gamma$



$D$

# 1-Липшиц функции



$$\forall x, y : |f(x) - f(y)| \leq |x - y| \Leftrightarrow \forall x : |f'(x)| \leq 1$$

$$WGAN = \operatorname{argmin}_G \operatorname{argmax}_{D \in 1\text{-lipschitz}} E_{x \in P_{real}(x)} D(x) - E_{x \in P_{fake}(x)} D(x)$$



# Способы сделать 1-липшиц регуляризацию

## WGAN

$$\begin{aligned}w &\leftarrow w + \alpha \cdot \text{RMSPProp}(w, g_w) \\w &\leftarrow \text{clip}(w, -c, c)\end{aligned}$$

---

**Algorithm 1** WGAN, our proposed algorithm. All experiments in the paper used the default values  $\alpha = 0.00005$ ,  $c = 0.01$ ,  $m = 64$ ,  $n_{\text{critic}} = 5$ .

---

**Require:** :  $\alpha$ , the learning rate.  $c$ , the clipping parameter.  $m$ , the batch size.  $n_{\text{critic}}$ , the number of iterations of the critic per generator iteration.

**Require:** :  $w_0$ , initial critic parameters.  $\theta_0$ , initial generator's parameters.

```
1: while  $\theta$  has not converged do
2:   for  $t = 0, \dots, n_{\text{critic}}$  do
3:     Sample  $\{x^{(i)}\}_{i=1}^m \sim \mathbb{P}_r$  a batch from the real data.
4:     Sample  $\{z^{(i)}\}_{i=1}^m \sim p(z)$  a batch of prior samples.
5:      $g_w \leftarrow \nabla_w [\frac{1}{m} \sum_{i=1}^m f_w(x^{(i)}) - \frac{1}{m} \sum_{i=1}^m f_w(g_\theta(z^{(i)}))]$ 
6:      $w \leftarrow w + \alpha \cdot \text{RMSPProp}(w, g_w)$ 
7:      $w \leftarrow \text{clip}(w, -c, c)$ 
8:   end for
9:   Sample  $\{z^{(i)}\}_{i=1}^m \sim p(z)$  a batch of prior samples.
10:   $g_\theta \leftarrow -\nabla_\theta \frac{1}{m} \sum_{i=1}^m f_w(g_\theta(z^{(i)}))$ 
11:   $\theta \leftarrow \theta - \alpha \cdot \text{RMSPProp}(\theta, g_\theta)$ 
12: end while
```

---

# WGAN-GP

---

**Algorithm 1** WGAN with gradient penalty. We use default values of  $\lambda = 10$ ,  $n_{\text{critic}} = 5$ ,  $\alpha = 0.0001$ ,  $\beta_1 = 0$ ,  $\beta_2 = 0.9$ .

---

**Require:** The gradient penalty coefficient  $\lambda$ , the number of critic iterations per generator iteration  $n_{\text{critic}}$ , the batch size  $m$ , Adam hyperparameters  $\alpha, \beta_1, \beta_2$ .


**Require:** initial critic parameters  $w_0$ , initial generator parameters  $\theta_0$ .

```

1: while  $\theta$  has not converged do
2:   for  $t = 1, \dots, n_{\text{critic}}$  do
3:     for  $i = 1, \dots, m$  do
4:       Sample real data  $\mathbf{x} \sim \mathbb{P}_r$ , latent variable  $\mathbf{z} \sim p(\mathbf{z})$ , a random number  $\epsilon \sim U[0, 1]$ .
5:        $\tilde{\mathbf{x}} \leftarrow G_\theta(\mathbf{z})$ 
6:        $\hat{\mathbf{x}} \leftarrow \epsilon \mathbf{x} + (1 - \epsilon) \tilde{\mathbf{x}}$ 
7:        $L^{(i)} \leftarrow D_w(\tilde{\mathbf{x}}) - D_w(\mathbf{x}) + \lambda (\|\nabla_{\hat{\mathbf{x}}} D_w(\hat{\mathbf{x}})\|_2 - 1)^2$ 
8:     end for
9:      $w \leftarrow \text{Adam}(\nabla_w \frac{1}{m} \sum_{i=1}^m L^{(i)}, w, \alpha, \beta_1, \beta_2)$ 
10:   end for
11:   Sample a batch of latent variables  $\{\mathbf{z}^{(i)}\}_{i=1}^m \sim p(\mathbf{z})$ .
12:    $\theta \leftarrow \text{Adam}(\nabla_\theta \frac{1}{m} \sum_{i=1}^m -D_w(G_\theta(\mathbf{z})), \theta, \alpha, \beta_1, \beta_2)$ 
13: end while

```

---



$$L = \underbrace{\mathbb{E}_{\tilde{\mathbf{x}} \sim \mathbb{P}_g} [D(\tilde{\mathbf{x}})] - \mathbb{E}_{\mathbf{x} \sim \mathbb{P}_r} [D(\mathbf{x})]}_{\text{Original critic loss}} + \lambda \underbrace{\mathbb{E}_{\hat{\mathbf{x}} \sim \mathbb{P}_{\hat{\mathbf{x}}}} [(\|\nabla_{\hat{\mathbf{x}}} D(\hat{\mathbf{x}})\|_2 - 1)^2]}_{\text{Our gradient penalty}}.$$

# Spectral normalisation

$$\|Ax\| \leq K\|x\|$$

$$\langle Ax, Ax \rangle \leq K^2 \langle x, x \rangle, \forall x \in I$$

$$\langle (A^T A - K^2)x, x \rangle \leq 0, \forall x \in I.$$

$$\begin{aligned} \langle (A^T A - K^2)x, x \rangle &= \langle (A^T A - K^2) \sum_i x_i v_i, \sum_j x_j v_j \rangle \\ &= \sum_i \sum_j x_i x_j \langle (A^T A - K^2)v_i, v_j \rangle \\ &= \sum_i (\lambda_i - K^2) x_i^2 \leq 0 \\ &\Rightarrow \sum_i (K^2 - \lambda_i) x_i^2 \geq 0. \end{aligned}$$

$$K^2 - \lambda_i \geq 0 \text{ for all } i = 1 \dots n.$$

$$\|f\|_{\text{Lip}} = \sup_x \sigma(\nabla f(x))$$

$$\nabla(g \circ f)(x) = \nabla g(f(x)) \nabla f(x).$$

$$\nabla(g \circ f)(x) = \nabla g(f(x)) \nabla f(x).$$

$$\sigma(\nabla f(x)) = \sup_{\|v\| \leq 1} \|[\nabla f(x)]v\|$$

$$\begin{aligned} \sigma(\nabla(g \circ f)(x)) &= \sup_{\|v\| \leq 1} \|[\nabla g(f(x))][\nabla f(x)]v\| \\ \sup_{\|v\| \leq 1} \|[\nabla g(f(x))][\nabla f(x)]v\| &\leq \sup_{\|u\| \leq 1} \|[\nabla g(f(x))]u\| \sup_{\|v\| \leq 1} \|[\nabla f(x)]v\|. \end{aligned}$$

$$\|g \circ f\|_{\text{Lip}} \leq \|g\|_{\text{Lip}} \|f\|_{\text{Lip}}.$$

# Power iteration

$$v_{t+1} = \frac{W^T W v_t}{\|W^T W v_t\|}$$

$$\begin{aligned} v_t &= \frac{(W^T W)^t \sum_i v_i e_i}{\|(W^T W)^t \sum_i v_i e_i\|} \\ &= \frac{\sum_i v_i \lambda_i^t e_i}{\|\sum_i v_i \lambda_i^t e_i\|} \\ &= \frac{v_1 \lambda_1^t \sum_i \frac{v_i}{v_1} \left(\frac{\lambda_i}{\lambda_1}\right)^t e_i}{\|v_1 \lambda_1^t \sum_i \frac{v_i}{v_1} \left(\frac{\lambda_i}{\lambda_1}\right)^t e_i\|}. \end{aligned}$$

$$\sigma(W) = \|Wv\| = u^T Wv.$$

# Progressive growing of GAN




Figure 1: Our training starts with both the generator (G) and discriminator (D) having a low spatial resolution of  $4 \times 4$  pixels. As the training advances, we incrementally add layers to G and D, thus increasing the spatial resolution of the generated images. All existing layers remain trainable throughout the process. Here  $[N \times N]$  refers to convolutional layers operating on  $N \times N$  spatial resolution. This allows stable synthesis in high resolutions and also speeds up training considerably. On the right we show six example images generated using progressive growing at  $1024 \times 1024$ .

Discriminator	Act.	Output shape	Params
Input image	—	$3 \times 1024 \times 1024$	—
Conv $1 \times 1$	LReLU	$16 \times 1024 \times 1024$	64
Conv $3 \times 3$	LReLU	$16 \times 1024 \times 1024$	2.3k
Conv $3 \times 3$	LReLU	$32 \times 1024 \times 1024$	4.6k
Downsample	—	$32 \times 512 \times 512$	—
Conv $3 \times 3$	LReLU	$32 \times 512 \times 512$	9.2k
Conv $3 \times 3$	LReLU	$64 \times 512 \times 512$	18k
Downsample	—	$64 \times 256 \times 256$	—
Conv $3 \times 3$	LReLU	$64 \times 256 \times 256$	37k
Conv $3 \times 3$	LReLU	$128 \times 256 \times 256$	74k
Downsample	—	$128 \times 128 \times 128$	—
Conv $3 \times 3$	LReLU	$128 \times 128 \times 128$	148k
Conv $3 \times 3$	LReLU	$256 \times 128 \times 128$	295k
Downsample	—	$256 \times 64 \times 64$	—
Conv $3 \times 3$	LReLU	$256 \times 64 \times 64$	590k
Conv $3 \times 3$	LReLU	$512 \times 64 \times 64$	1.2M
Downsample	—	$512 \times 32 \times 32$	—
Conv $3 \times 3$	LReLU	$512 \times 32 \times 32$	2.4M
Conv $3 \times 3$	LReLU	$512 \times 32 \times 32$	2.4M
Downsample	—	$512 \times 16 \times 16$	—
Conv $3 \times 3$	LReLU	$512 \times 16 \times 16$	2.4M
Conv $3 \times 3$	LReLU	$512 \times 16 \times 16$	2.4M
Downsample	—	$512 \times 8 \times 8$	—
Conv $3 \times 3$	LReLU	$512 \times 8 \times 8$	2.4M
Conv $3 \times 3$	LReLU	$512 \times 8 \times 8$	2.4M
Downsample	—	$512 \times 4 \times 4$	—
Minibatch stddev	—	$513 \times 4 \times 4$	—
Conv $3 \times 3$	LReLU	$512 \times 4 \times 4$	2.4M
Conv $4 \times 4$	LReLU	$512 \times 1 \times 1$	4.2M
Fully-connected	linear	$1 \times 1 \times 1$	513
Total trainable parameters			<b>23.1M</b>

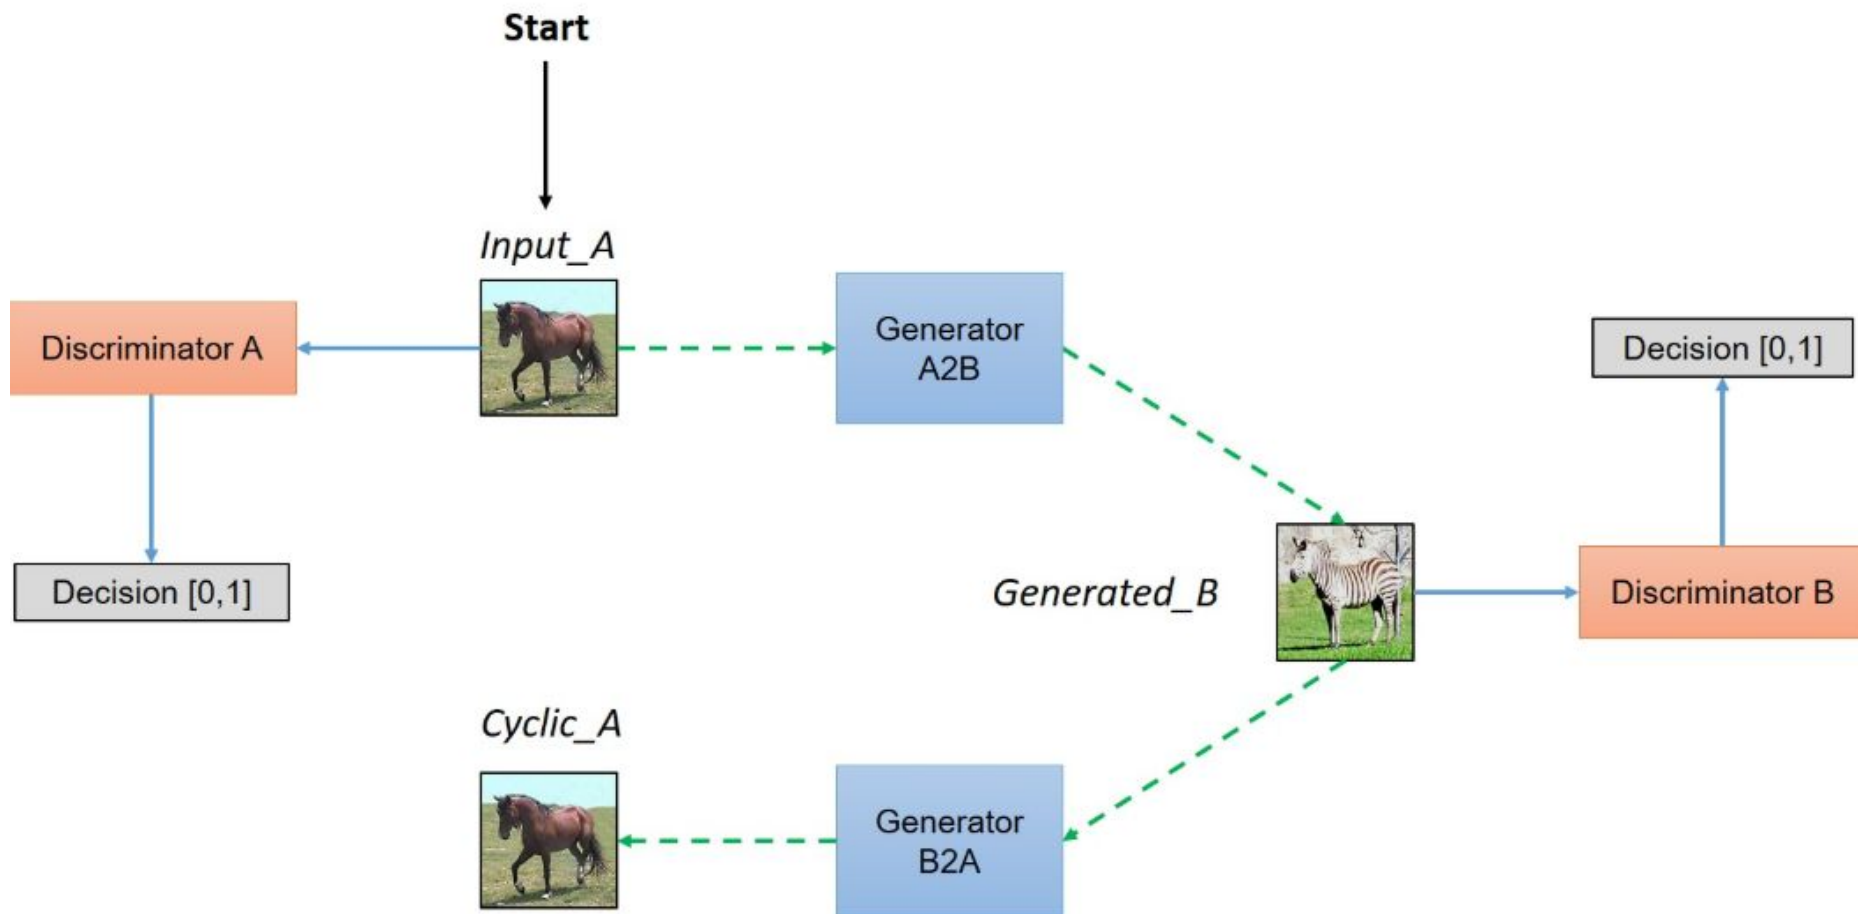


# Вопросы

1. Почему wasserstein gan обучается лучше и стабильнее обычного (проблема kl divergence)
2. Что такое расстояние тракториста (earth mover distance), как его можно считать (ЛП, разница функций распределений по модулю в одномерном случае)
3. Два способа из 3 поддерживать 1 липшиц св-во, интуиция
4. Чем помогает progressive growing of gans, как он устроен

# Cycle GAN





# Источники

<https://www.youtube.com/channel/UC2ggituuWvvrHHHiaDH1dIQ>

<https://blog.acolyer.org/2018/05/10/progressive-growing-of-gans-for-improved-quality-stability-and-variation/>

[https://medium.com/@jonathan\\_hui/gan-wasserstein-gan-wgan-gp-6a1a2aa1b490](https://medium.com/@jonathan_hui/gan-wasserstein-gan-wgan-gp-6a1a2aa1b490)

[https://sthalles.github.io/advanced\\_gans/](https://sthalles.github.io/advanced_gans/)

<https://christiancosgrove.com/blog/2018/01/04/spectral-normalization-explained.html>

<https://github.com/soumith/ganhacks>