



НАЦИОНАЛЬНЫЙ ИССЛЕДОВАТЕЛЬСКИЙ
УНИВЕРСИТЕТ

БПМИ-191: Абрамов Арсений

MULTIPLYING MATRICES WITHOUT MULTIPLYING

Москва, 2022



- Матричное умножение лежит в основе многих алгоритмов и возникает почти всюду в ТИ и МО.
- Честное матричное умножение имеет кубическую сложность (плохо).
- Approximating Matrices Multiplication (АММ) ищет баланс между точностью и эффективностью по времени / памяти.
- Активно разрабатываются новые методы АММ, обладающие ещё большей эффективностью. Много методов АММ работают при допущениях.



Линейный оператор

$$\mathbf{A} \in R^{N \times D}, \mathbf{B} \in R^{D \times M}, N \gg D \geq M$$

Data

$$\mathbf{V}_A, \mathbf{V}_B \in R^{D \times d}, d \ll D$$

sparse

$$\mathbf{AB} \sim (\mathbf{AV}_A)(\mathbf{V}_B^T \mathbf{B})$$

(Если коротко, то проблема сводится к честному произведению в меньшей размерности через линейные преобразования \mathbf{A} и \mathbf{B} .)



Assumptions:

- Matrices are tall
- Matrices are relatively dense
- Resident in a single machine's memory
- **Exists a training set \tilde{A} .**

Given a matrix $A \in \mathbb{R}^{m \times n}$, let us define:

- A is a **fat matrix** if $m \leq n$ and $\text{null}(A^T) = \{0\}$
- A is a **tall matrix** if $m \geq n$ and $\text{range}(A) = \mathbb{R}^n$

ядро

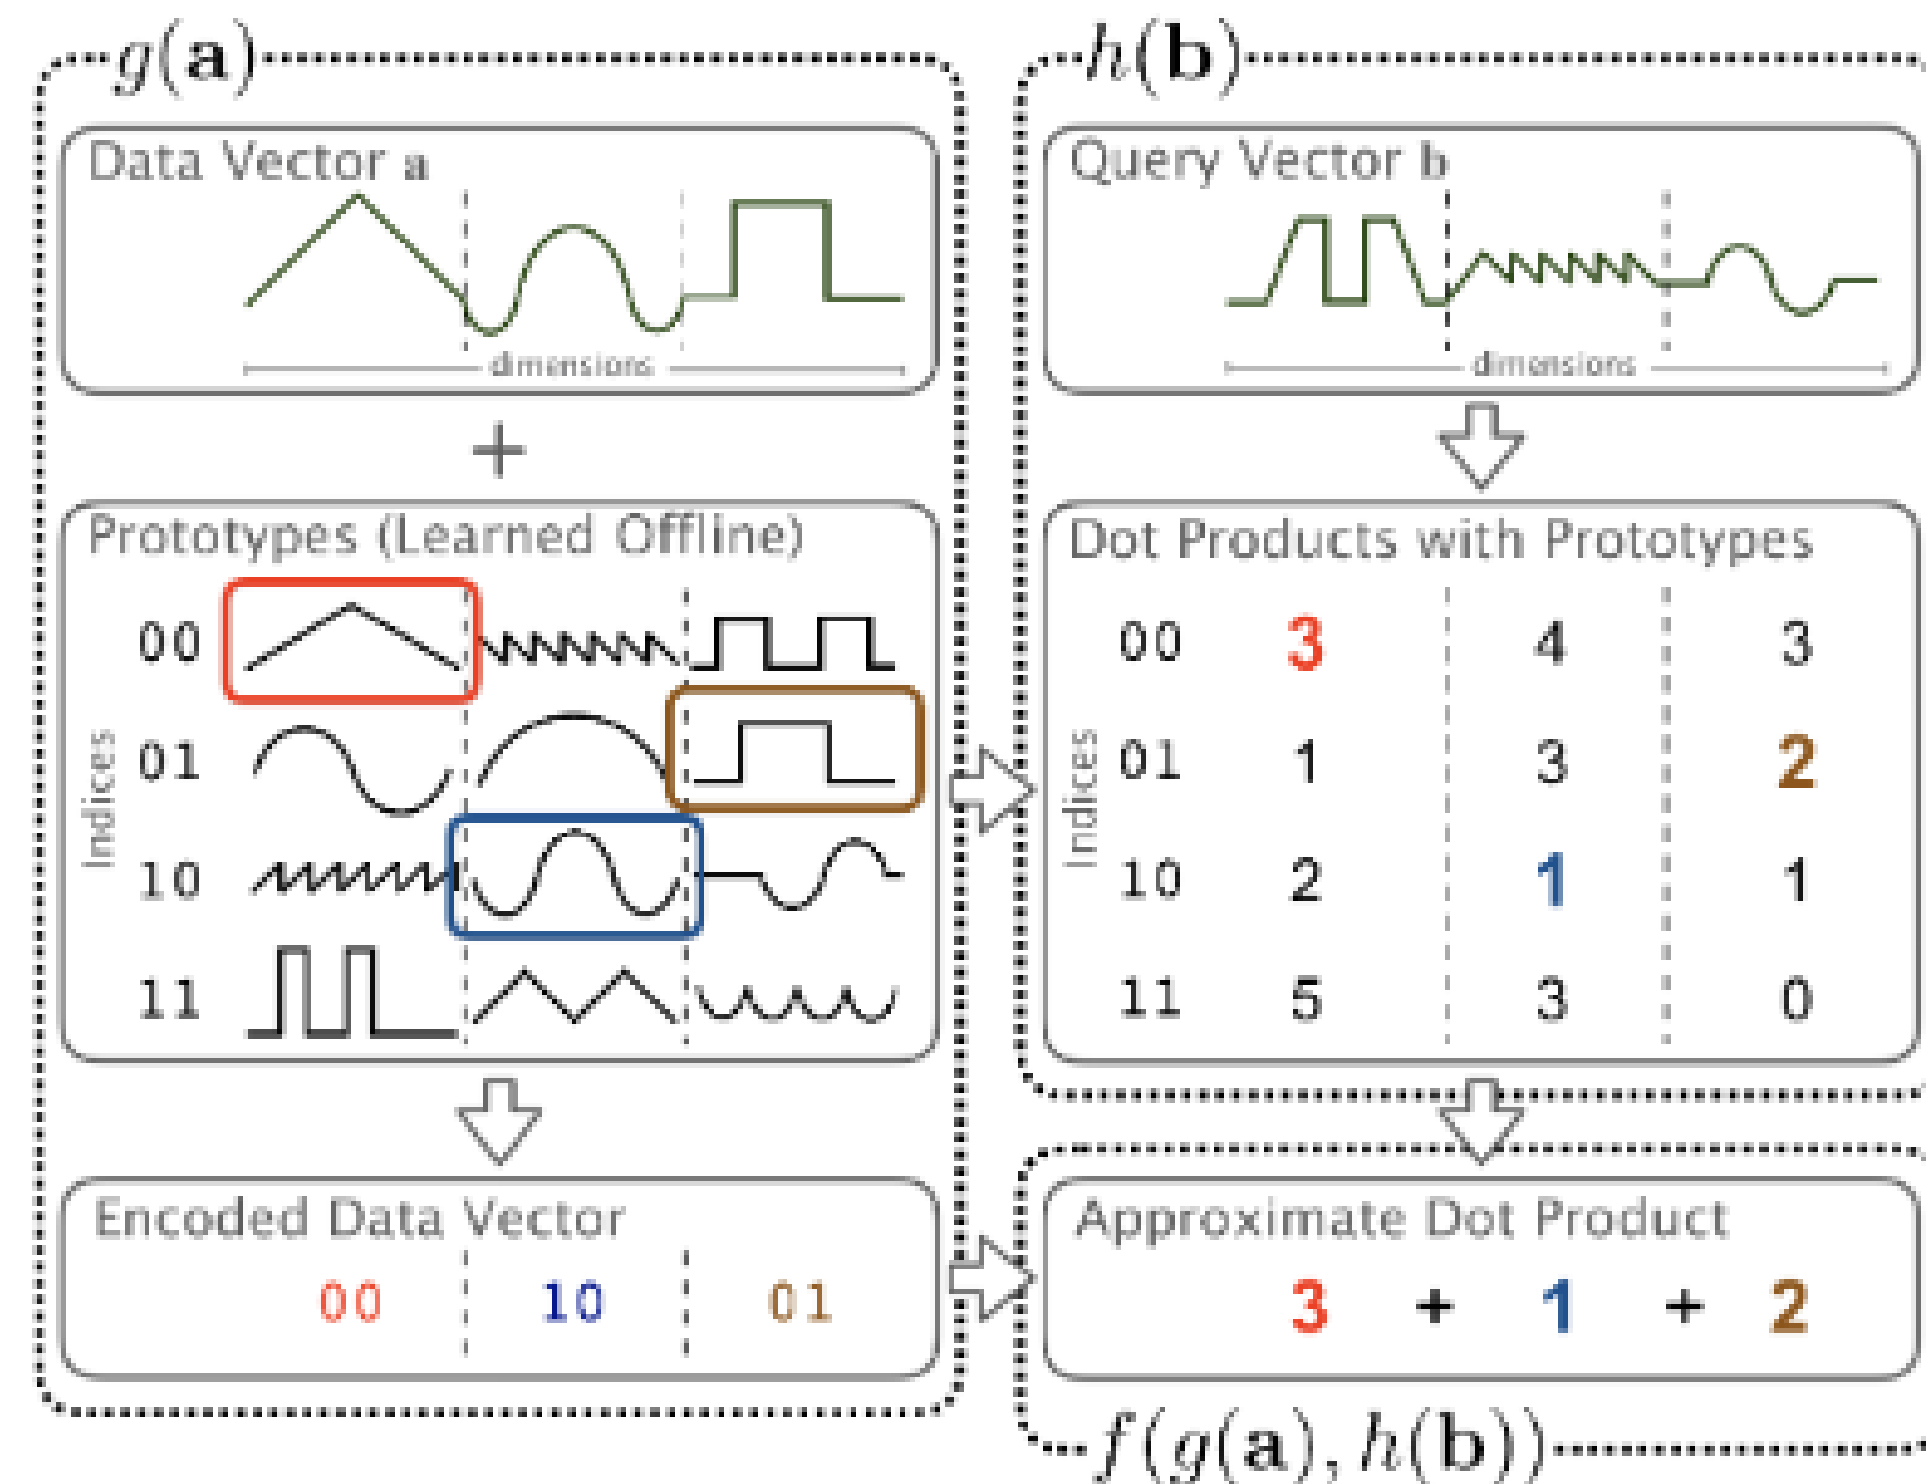
образ



MADDNESS employs a nonlinear preprocessing function and reduces the problem to table lookups.

If B is *known ahead of time*, MADDNESS **requires no multiply-add operations.**

Product Quantization



1. **Prototype Learning** - In an initial, offline training phase, cluster the rows of \mathbf{A} (or a training set $\tilde{\mathbf{A}}$) using K-means to create prototypes. A separate K-means is run in each of C disjoint subspaces to produce C sets of K prototypes.
2. **Encoding Function, $g(\mathbf{a})$** - Determine the most similar prototype to \mathbf{a} in each subspace. Store these assignments as integer indices using $C \log_2(K)$ bits.
3. **Table Construction, $h(\mathbf{B})$** - Precompute the dot products between \mathbf{b} and each prototype in each subspace. Store these partial dot products in C lookup tables of size K .
4. **Aggregation, $f(\cdot, \cdot)$** - Use the indices and tables to lookup the estimated partial $\mathbf{a}^\top \mathbf{b}$ in each subspace, then sum the results across all C subspaces.



Hash function family

Algorithm 1 MADDNESSHASH

```
1: Input: vector  $\mathbf{x}$ , split indices  $j^1, \dots, j^4$ , split thresh-  
   olds  $v^1, \dots, v^4$   
2:  $i \leftarrow 1$  // node index within level of tree  
3: for  $t \leftarrow 1$  to 4 do  
4:    $v \leftarrow v_i^t$  // lookup split threshold for node  $i$  at level  $t$   
5:    $b \leftarrow x_{j^t} \geq v ? 1 : 0$  // above split threshold?  
6:    $i \leftarrow 2i - 1 + b$  // assign to left or right child  
7: end for  
8: return  $i$ 
```

Learning the hash function parameters

Algorithm 2 Adding The Next Level to the Hashing Tree

```
1: Input: buckets  $\mathcal{B}_1^{t-1}, \dots, \mathcal{B}_{2^{t-1}}^{t-1}$ , training matrix  $\tilde{A}$   
   // greedily choose next split index and thresholds  
2:  $\hat{\mathcal{J}} \leftarrow \text{heuristic\_select\_idxs}(\mathcal{B}_1^{t-1}, \dots, \mathcal{B}_{2^{t-1}}^{t-1})$   
3:  $l^{\min}, j^{\min}, v^{\min} \leftarrow \infty, \text{NaN}, \text{NaN}$   
4: for  $j \in \hat{\mathcal{J}}$  do  
5:    $l \leftarrow 0$  // initialize loss for this index to 0  
6:    $v \leftarrow []$  // empty list of split thresholds  
7:   for  $i \leftarrow 1$  to  $2^{t-1}$  do  
8:      $v_i, l_i \leftarrow \text{optimal\_split\_threshold}(j, \mathcal{B}_i^{t-1})$   
9:      $\text{append}(v, v_i)$  // append threshold for bucket  $i$   
10:     $l \leftarrow l + l_i$  // accumulate loss from bucket  $i$   
11:   end for  
12:   if  $l < l^{\min}$  then  
13:      $l^{\min} \leftarrow l, j^{\min} \leftarrow j, v^{\min} \leftarrow v$  // new best split  
14:   end if  
15: end for  
   // create new buckets using chosen split  
16:  $\mathcal{B} \leftarrow []$   
17: for  $i \leftarrow 1$  to  $2^{t-1}$  do  
18:    $\mathcal{B}_{\text{below}}, \mathcal{B}_{\text{above}} \leftarrow \text{apply\_split}(v_i^{\min}, \mathcal{B}_i^{t-1})$   
19:    $\text{append}(\mathcal{B}, \mathcal{B}_{\text{below}})$   
20:    $\text{append}(\mathcal{B}, \mathcal{B}_{\text{above}})$   
21: end for  
22: return  $\mathcal{B}, l^{\min}, j^{\min}, v^{\min}$ 
```

$$\mathcal{L}(j, \mathcal{B}) \triangleq \sum_{\mathbf{x} \in \mathcal{B}} \left(x_j - \frac{1}{|\mathcal{B}|} \sum_{\mathbf{x}' \in \mathcal{B}} x'_j \right)^2$$

$$\mathcal{L}(\mathcal{B}) \triangleq \sum_j \mathcal{L}(j, \mathcal{B}).$$



Optimizing prototypes

$$\mathbf{P} \triangleq (\mathbf{G}^\top \mathbf{G} + \lambda \mathbf{I})^{-1} \mathbf{G}^\top \tilde{\mathbf{A}}.$$

$$\mathbf{P} \in \mathbb{R}^{K \times C \times D}$$

Param = 1

$$\tilde{\mathbf{A}} \approx \bar{\mathbf{G}} \mathbf{P}.$$




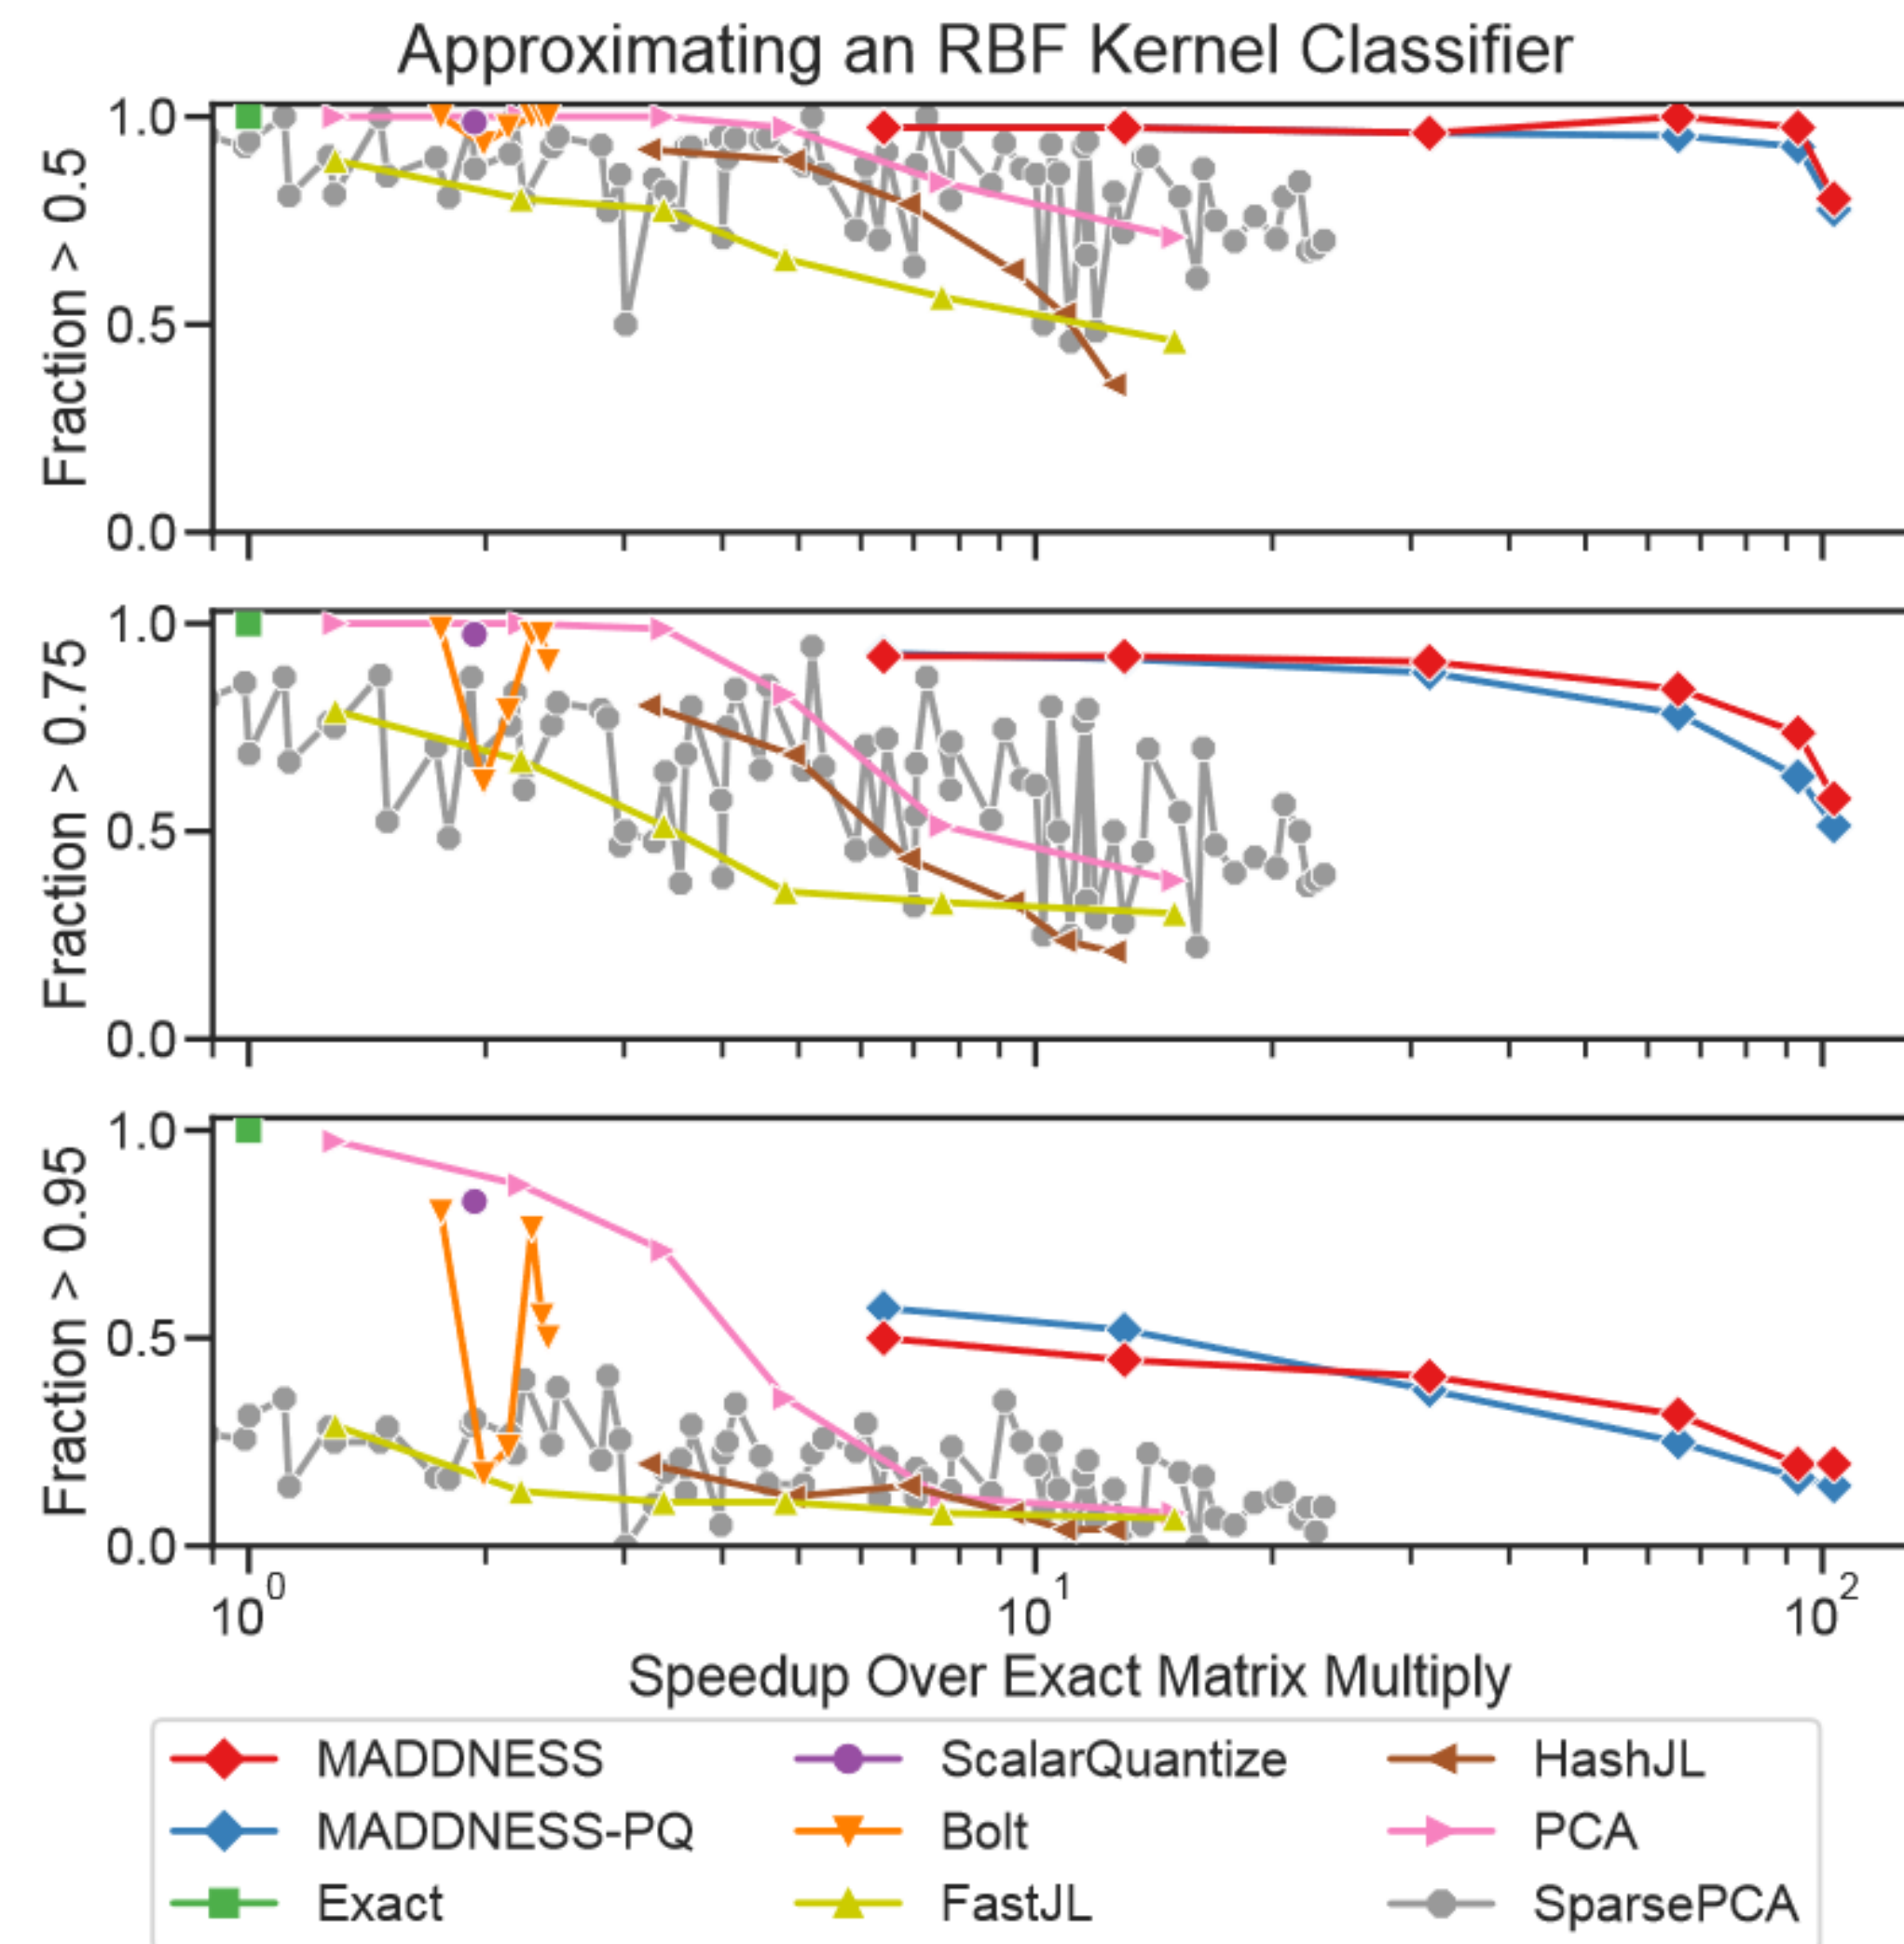

Fast 8-bit aggregation

Let $\mathbf{T} \in \mathbb{R}^{M \times C \times K}$ be the tensor of lookup tables for all M columns of \mathbf{B} . Given the encodings \mathbf{G} , the function $f(\cdot, \cdot)$ is defined as

$$f(g(\mathbf{A}), h(\mathbf{B}))_{n,m} \triangleq \sum_{c=1}^C \mathbf{T}_{m,c,k}, \quad k = g^{(c)}(\mathbf{a}_n). \quad (9)$$

Addition \longrightarrow Averaging







$$AB \approx (AV_A)(V_B^T B).$$



Learning Hash Function



<https://smarturl.it/Maddness>



ИСТОЧНИКИ

(Источник)

ИСТОЧНИК

Источник

- *Davis Blalock, John Guttag: Multiplying Matrices Without Multiplying*
URL: <https://arxiv.org/pdf/2106.10860.pdf>



НАЦИОНАЛЬНЫЙ ИССЛЕДОВАТЕЛЬСКИЙ
УНИВЕРСИТЕТ