

Mastering the game of Go with deep neural networks

David Silver, Aja Huang, Christopher Maddison, 2016, Nature

Yakshimamedov Vladimir, 192

What is Go?

Two-player, zero-sum,
complete information.

States(positions) \approx

$2.081681994 * 10^{170}$

Chess positions upper bound -
 10^{123}





Policy and value functions

Policy - f : state \rightarrow probabilities to choose each action.


Softmax can be added to choose one action instead.

Value function - $v_{\pi}(s) = \mathbb{E}_{\pi}[G_t | S_t = s]$

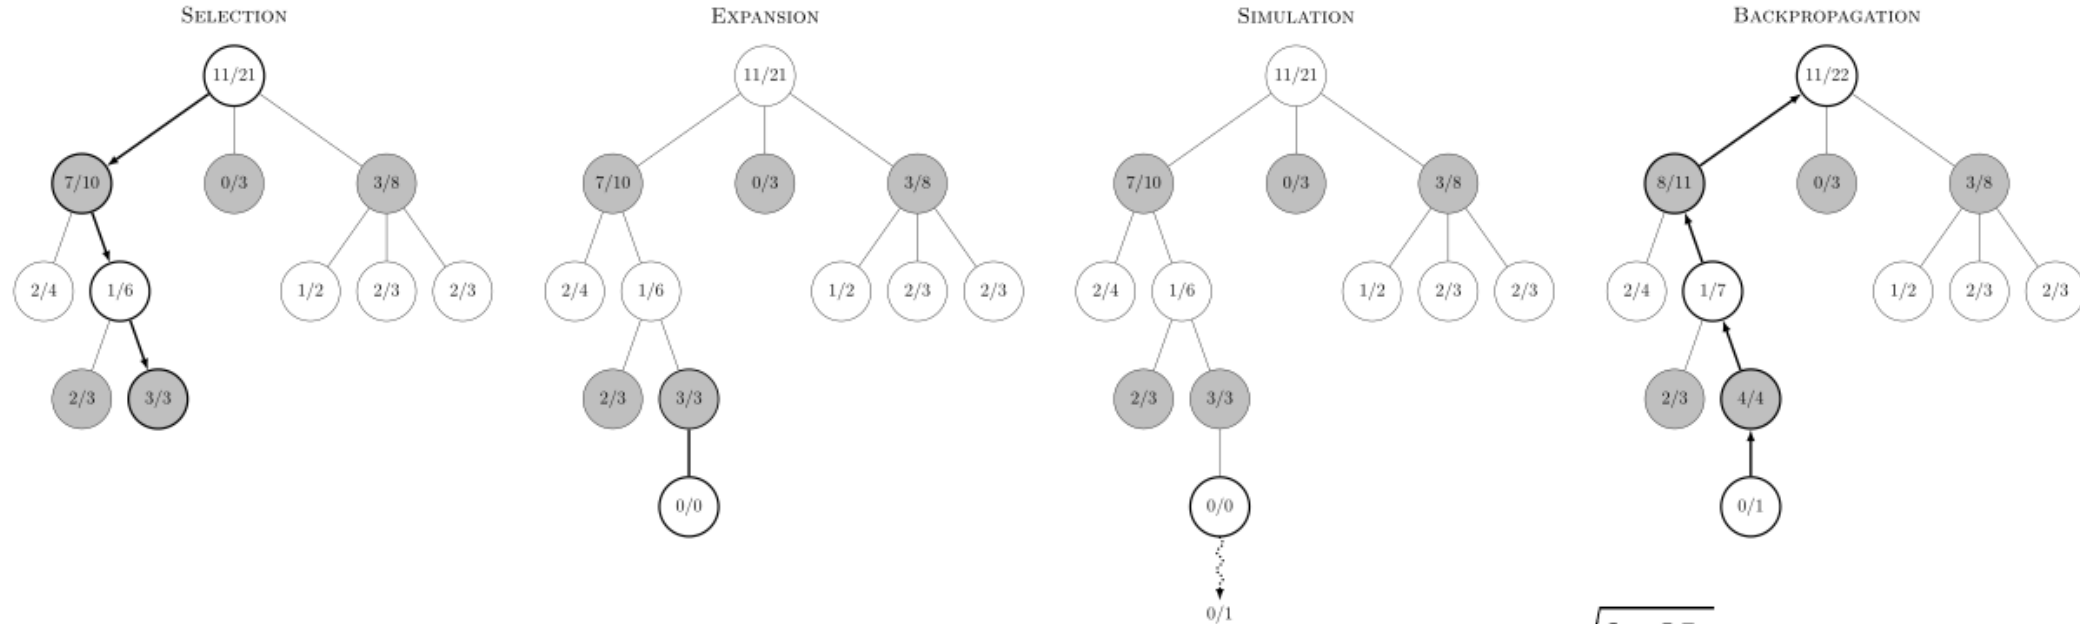
Action-value function - $q_{\pi}(s, a) = \mathbb{E}_{\pi}[G_t | S_t = s, A_t = a]$



Full AlphaGo structure



Monte-Carlo Tree Search (MCTS)



Selection – Choose action: $a_t = \underset{a}{\operatorname{argmax}} (Q(s_t, a) + u(s_t, a))$

$$\frac{w_i}{n_i} + c \sqrt{\frac{\ln N_i}{n_i}}$$

Expansion – If not terminal, add children.

Simulation – Choose a child and play a complete (random or simple policy) payout.

Backpropagation – Update win count and play count on parents.

Chad Rollout Policy

Chooses an action before you ask. (2 μ s to choose an action)

Trained on like 6 features.

Only 32 filters of conv layers

What are symmetries?

Accuracy on test set: 24.2%

Virgin Supervised Learning Policy

Needs 3 ms to choose

Spends its life looking at 50 features.

256 filters of convolutional layers

Spins the board for no reason.

(Mean of 8 evaluations – flipping the colour and rotating the board)

Took 3 weeks to train.

Accuracy on test set: 57%

RL policy network

Initialization:

Structure and weights are the same as SL policy.

Step:

Play a batch of games with one of the previous iterations

Update weights using stochastic gradient.

$$\Delta \rho \propto \frac{\partial \log p_{\rho}(a_t | s_t)}{\partial \rho} z_t.$$

(log likelihood)

RL value network

Initialization:

Structure is the same as SL policy, but with 1 output.

Step:

Generate a lot of games.

Choose one position from every game and predict the outcome of that game.

Subsequent positions from fewer games are significantly worse.

Optimizing MSE

Monte-Carlo in AlphaZero

Selection:

$$a_t = \underset{a}{\operatorname{argmax}} (Q(s_t, a) + u(s_t, a))$$

$P(s, a)$ - prior probability to choose this action from SL policy network

Exploration function - $u(s, a) = c_{\text{puct}} P(s, a) \frac{\sqrt{\sum_b N_r(s, b)}}{1 + N_r(s, a)}$

Expansion:

1. Reach the threshold for playout counts.
2. Creating node with all action edges:

For every edge:

$N_r(s, a)$ - count of playout simulations for every edge

$N_v(s, a)$ - count of value network evaluations for the $s' = s + a$.

$W_r(s, a)$ - wins over $N_r(s, a)$ playouts.

$W_v(s, a)$ - value sum over $N_v(s, a)$ evaluations.

$Q(s, a)$ - state-action evaluation.

Initialized as 0.

Monte-Carlo in AlphaZero (Part 2)

Simulation (Evaluation):

Evaluate new state with value network
if it wasn't already.

Playout this position using rollout policy.

Backup:

For each edge:

$$N_r(s, a) += 1$$

$$N_v(s, a) += 1$$

$$W_r(s, a) += 0 \text{ or } 1$$

$W_v(s, a) +=$ value network evaluation of
the new state

$$Q(s, a) = (1 - \lambda) \frac{W_v(s, a)}{N_v(s, a)} + \lambda \frac{W_r(s, a)}{N_r(s, a)}.$$


Full AlphaGo structure

Rollout – fast policy for
playouts in MCTS.

SL policy – pre-train for RL
policy + prior probabilities for
MCTS.

RL policy – better SL for
generating games

Value network – finally
something evaluating
positions.



Parameters

Lambda = 0.5 is the optimal parameter.

But even relying only on value network or only MCTS gives decent results.

Hell, even MCTS on rollouts is probably better than me.

| Short name | Policy network | Value network | Rollouts | Mixing constant | Policy GPUs | Value GPUs | Elo rating |
|----------------|----------------|---------------|----------|-----------------|-------------|------------|------------|
| α_{rvp} | p_σ | v_θ | p_π | $\lambda = 0.5$ | 2 | 6 | 2890 |
| α_{vp} | p_σ | v_θ | — | $\lambda = 0$ | 2 | 6 | 2177 |
| α_{rp} | p_σ | — | p_π | $\lambda = 1$ | 8 | 0 | 2416 |
| α_{rv} | $[p_\tau]$ | v_θ | p_π | $\lambda = 0.5$ | 0 | 8 | 2077 |
| α_v | $[p_\tau]$ | v_θ | — | $\lambda = 0$ | 0 | 8 | 1655 |
| α_r | $[p_\tau]$ | — | p_π | $\lambda = 1$ | 0 | 0 | 1457 |
| α_p | p_σ | — | — | — | 0 | 0 | 1517 |

$[p_\tau]$ - Monte-Carlo Tree Search on only Rollout policy

p_σ - Supervised Learning policy

Distribution (or how to throw 176 GPUS at a model)

Monte Carlo Changes – add constant to rollout count at the start of an iteration to discourage other threads to follow the same path.

CPU – rollouts.

GPU – value and policy networks.

| Short name | Computer Player | Version | Time settings | CPU | GPU | KGS Rank | Elo |
|------------------------|----------------------------|-------------------|---------------|------|-----|----------|------|
| α_{rvp}^d | Distributed <i>AlphaGo</i> | See Methods | 5 seconds | 1202 | 176 | – | 3140 |
| α_{rvp} | <i>AlphaGo</i> | See Methods | 5 seconds | 48 | 8 | – | 2890 |
| <i>CS</i> | CrazyStone | 2015 | 5 seconds | 32 | – | 6d | 1929 |
| <i>ZN</i> | Zen | 5 | 5 seconds | 8 | – | 6d | 1888 |
| <i>PC</i> | Pachi | 10.99 | 400,000 sims | 16 | – | 2d | 1298 |
| <i>FG</i> | Fuego | svn1989 | 100,000 sims | 16 | – | – | 1148 |
| <i>GG</i> | GnuGo | 3.8 | level 10 | 1 | – | 5k | 431 |
| <i>CS</i> ₄ | CrazyStone | 4 handicap stones | 5 seconds | 32 | – | – | 2526 |
| <i>ZN</i> ₄ | Zen | 4 handicap stones | 5 seconds | 8 | – | – | 2413 |
| <i>PC</i> ₄ | Pachi | 4 handicap stones | 400,000 sims | 16 | – | – | 1756 |